

# Certified Kubernetes Application Developer (CKAD)





# Ken Sipe

Distribute Application Engineer

Apache Mesos Contributor  
Apache Committer Myriad, Open DCOS  
Developer: Embedded, C++, Java, Groovy, Grails, C#, GoLang

 @KenSipe

---

# Agenda

- CKAD - Expectation
- Setting Up
- Core Concepts
- Configuration
- Pod Design
- Services / Networking
- Persistence



---

# CKAD Expectations

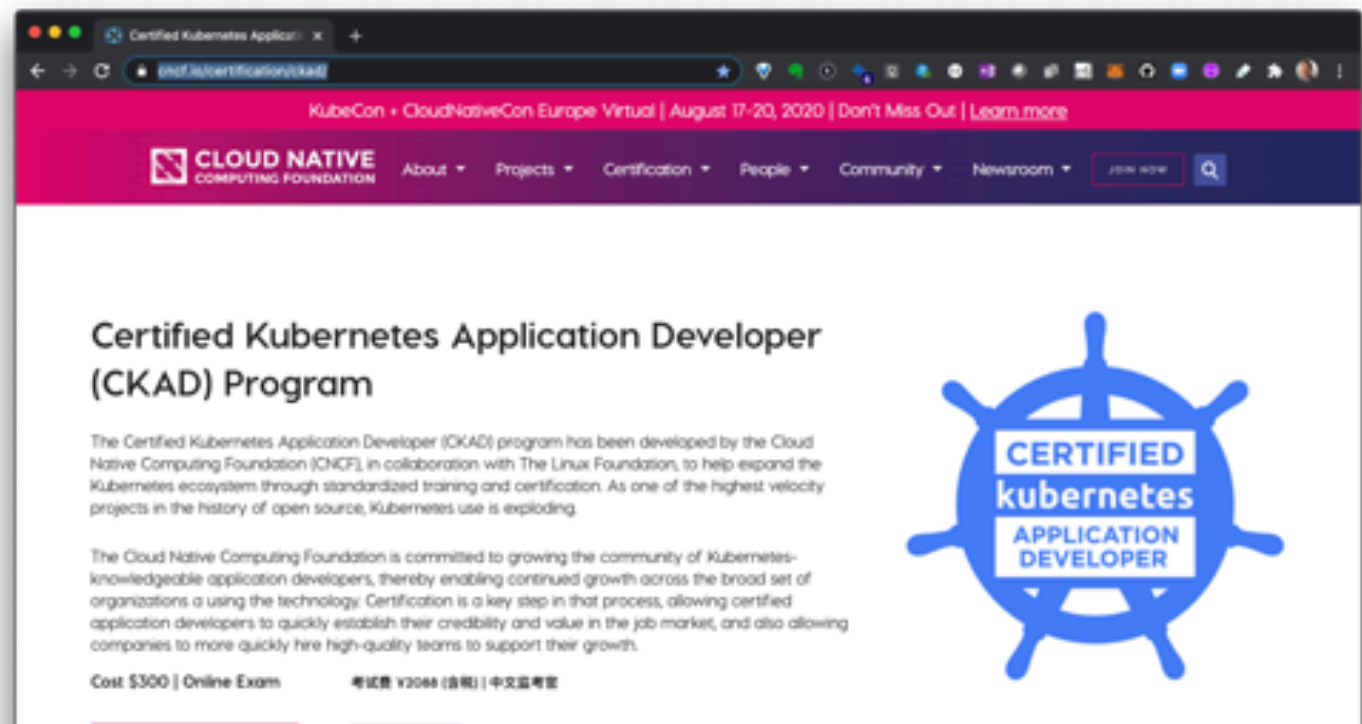
---

## Expectations

- Number of Questions
- Time per Question
- Desk / Proxy
- Cost / 2x
- Topics / percents
- Signing up (links)

Sign up

<https://www.cncf.io/certification/ckad/>



## Cost

- \$300
  - Includes 1 Retest
- 2 hours



---

## Time

- 2 hours
- 20 questions (weighted)
- **6 mins / Question**





## Taking Test

- Remote Proxy
  - Video scan desk area
  - Closed uninterruptible space
  - No-one in the area
  - Microphone, Video and full view of all screens
  - Clear Desktop
- Reliable Internet



---

## Browser

- Current version of Chrome or Chromium
- Limited Access
  - Test
  - [kubernetes.io](https://kubernetes.io)
  - No Google!
- Search is on [kubernetes.io](https://kubernetes.io)

## Browser

- Know your weakness

Examples using `kubectl` and JSONPath expressions:

```
ctl get pods -o json
ctl get pods -o=jsonpath='{@}'
ctl get pods -o=jsonpath='{.items[0]}'
ctl get pods -o=jsonpath='{.items[0].metadata.name}'
ctl get pods -o=jsonpath="{.items[*]['metadata.name', 'status.capacity']}"
ctl get pods -o=jsonpath='{range .items[*]}{.metadata.name}{"\t"}{.status.s'
```

<https://kubernetes.io/docs/reference/kubectl/jsonpath/>

<https://kubernetes.io/docs/reference/kubectl/cheatsheet/>

## 13% - Core Concepts

- Understand Kubernetes API primitives
- Create and configure basic Pods

## 8% - State Persistence

- Understand PersistentVolumeClaims for storage

## 20% - Pod Design

- Understand Deployments and how to perform rolling updates
- Understand Deployments and how to perform rollbacks
- Understand Jobs and CronJobs
- Understand how to use Labels, Selectors, and Annotations

## 18% - Observability

- Understand LivenessProbes and ReadinessProbes
- Understand container logging
- Understand how to monitor applications in Kubernetes

## 18% - Configuration

- Understand ConfigMaps
- Understand SecurityContexts
- Define an application's resource requirements
- Create & consume Secrets
- Understand ServiceAccounts

## 13% - Services & Networking

- Understand Services
- Demonstrate basic understanding of NetworkPolicies

## 13% - Services & Networking

- Understand Services
- Demonstrate basic understanding of NetworkPolicies

## 10% Multi-Container Pods

- Understand Multi-Container Pod design patterns (e.g. ambassador, adapter, sidecar)

Primary ID  
(non-expired and including photograph and signature)

Secondary ID  
(non-expired and including signature with Candidate name in Latin characters)

Passport

Government-issued driver's license/permit

Government-Issued local language ID (with photo and signature)

National Identity card

State or province-issued identity card

Debit (ATM) Card

Credit Card

Health Insurance Card

U.S. Social Security Card



<https://docs.linuxfoundation.org/tc-docs/certification/lf-candidate-handbook/candidate-requirements>

© 2020 D2iQ, Inc. All Rights Reserved.

---

## Passing Test

66 / 100



## Sign up and Schedule



The screenshot shows the login page of the Linux Foundation training portal. At the top, there is a dark blue banner with the Linux Foundation logo on the left and a network diagram with yellow nodes and blue lines on the right. Below the banner, the page has a white background. The heading "Sign In" is centered in blue. Below it, a link "New user? Create an account" is displayed in a smaller font. There are two input fields: the first is labeled "username" and the second is a password field with masked characters. Below the password field is a link "Forgot Password?". A blue "Sign In" button is centered below the links. Underneath the button, the text "Or Sign In with" is followed by four social media icons: Google, GitHub, LinkedIn, and Facebook. At the bottom right of the page, there is a small "Artisan" logo.

<https://trainingportal.linuxfoundation.org/>

© 2020 D2iQ, Inc. All Rights Reserved.



---

# CKAD Setting up



---

## Expectations

- VI skills and setup
- Shell in browser
- Browser setup and bookmarks
- Time savers

# VI

- Shell inside Browser
- alias vi=vim

```
cat ~/.vimrc
:set number
:set et
:set sw=2 ts=2 sts=2
```

- “number” - sets left col line nums
- “et” - expand tab
- “sw” - shiftwidth
- “ts” - tab stop
- “sts” - softtabstop

\*\* configuration for 2 spaces for a tab

## VI keys

- Useful keys:

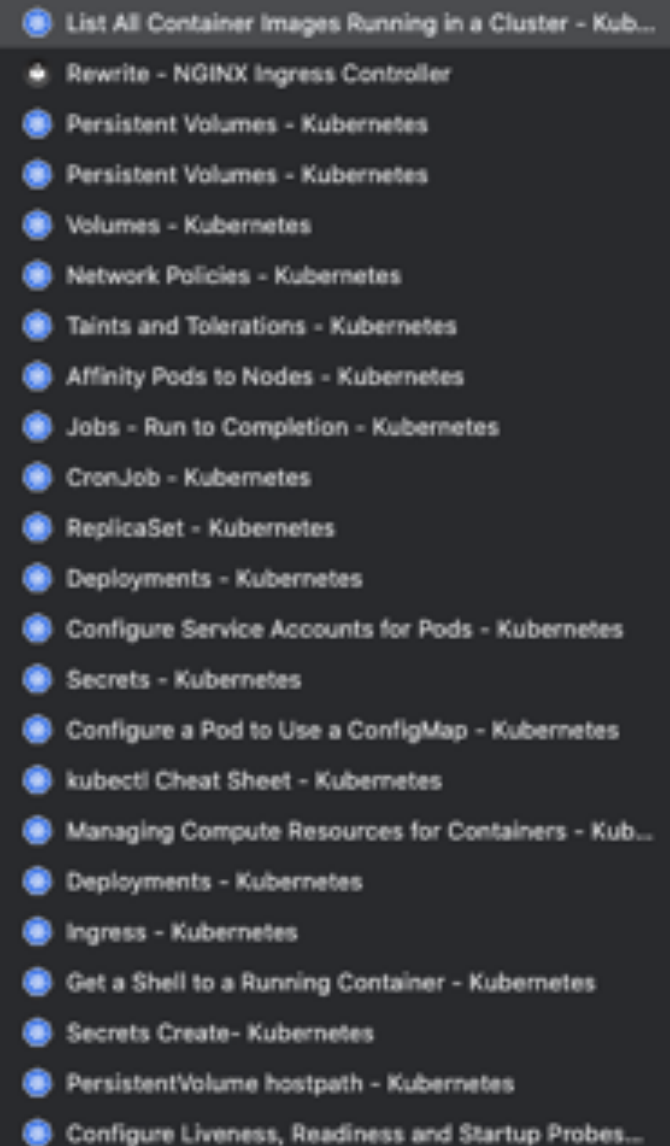
- ``/` - search
- `i` - insert / edit mode
- `<esc>` exit edit mode
- `:`` command prompt
- `:q!` - quit without writing
- `:wq` - write and quit
- ``A`` - append to end of line

- `:undo`` - undo
- ``dd`` - delete line
- ``cw`` - change word
- ``dw`` - delete word

<https://medium.com/free-code-camp/how-not-to-be-afraid-of-vim-anymore-ec0b7264b0ae>

## Browser

- Browser bookmarks
- 2 browser windows open
  - Test
  - Docs



A screenshot of a browser's bookmark bar, displaying 25 bookmarks related to Kubernetes. Each bookmark is preceded by a blue star icon. The bookmarks are as follows:

- List All Container Images Running in a Cluster - Kub...
- Rewrite - NGINX Ingress Controller
- Persistent Volumes - Kubernetes
- Persistent Volumes - Kubernetes
- Volumes - Kubernetes
- Network Policies - Kubernetes
- Taints and Tolerations - Kubernetes
- Affinity Pods to Nodes - Kubernetes
- Jobs - Run to Completion - Kubernetes
- CronJob - Kubernetes
- ReplicaSet - Kubernetes
- Deployments - Kubernetes
- Configure Service Accounts for Pods - Kubernetes
- Secrets - Kubernetes
- Configure a Pod to Use a ConfigMap - Kubernetes
- kubectl Cheat Sheet - Kubernetes
- Managing Compute Resources for Containers - Kub...
- Deployments - Kubernetes
- Ingress - Kubernetes
- Get a Shell to a Running Container - Kubernetes
- Secrets Create- Kubernetes
- PersistentVolume hostpath - Kubernetes
- Configure Liveness, Readiness and Startup Probes...

# Aliases

- Required

```
alias k=kubectl
```

- Nice to have

```
alias kdr="kubectl --dry-run -o=yaml "  
alias kget="kubectl -o=yaml get "  
alias desc="kubectl describe "
```

---

## Unix 1 Liners

```
while true; do date >> /var/log/app.txt; sleep 5; done
```

```
i=0; while true; do echo "$i: $(date)"; i=$((i+1)); sleep 1; done
```

```
mkdir -p collect; while true; do cat /var/data/* > /collect/data.text; sleep 10; done
```

```
a=10;b=5; if [ $a -le $b ]; then echo "a is small"; else echo "b is small"; fi
```

---

## Unix Skills

- `ctrl+r` search for previous command
- `<esc> .` - last argument used
  - (useful when you just edited yaml, now want to apply)

## Lab: Setup Kind

- Install [Kind](#)

```
kensipe@kens-mbp-2: ~/presentations/2020/ckad/labs
$ kind create cluster
Creating cluster "kind" ...
  ✓ Ensuring node image (kindest/node:v1.17.8) 📜
  ✓ Preparing nodes 📦
  ✓ Writing configuration 📄
  ✓ Starting control-plane 🚦
  ✓ Installing CNI 🌐
  ✓ Installing StorageClass 💾
Set kubectl context to "kind-kind"
You can now use your cluster with:

kubectl cluster-info --context kind-kind

Have a nice day! 🌞
$ kind create cluster --config=./kind.yaml
```



## K8S Manifest Shortcuts

```
k run nginx --image=nginx (deploy)
```

```
k run nginx --image=nginx --restart=Never (pod)
```

```
k run nginx --image=nginx --restart=OnFailure (job)
```

```
k run nginx --image=nginx --restart=OnFailure --schedule="* * * * *" (cronjob)
```

**\*\* pre- 1.18**

## K8S Manifest Shortcuts

```
k create deployment nginx --image=nginx (deploy)
```

```
k run nginx --image=nginx --restart=Never (pod)
```

```
k create job nginx --image=nginx (job)
```

```
k create cronjob nginx --image=nginx --schedule="* * * * *" (cronjob)
```

**\*\* post- 1.18**

# Rapid Deployment Creation

```
k create deployment nginx --image=nginx --dry-run -o yaml > test-1.yaml
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: nginx
  name: nginx
spec:
  replicas: 1
  selector:
    matchLabels:
      app: nginx
  strategy: {}
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: nginx
    spec:
      containers:
        - image: nginx
          name: nginx
```

---

## Know Explain

- k explain pod.spec
- k explain pod --recursive
- k explain cronjob.spec.jobTemplate --recursive

- need details on nodeselector

```
k explain pod.spec | grep -i nodeselector
```

- Know grep

```
k explain pod.spec | grep -i -C 4 nodeselector
```

---

## Know Manifest Layout

- Can't remember?
  - Dry-run output a pod
  - Or explain pod

`apiVersion:`

`kind:`

`metadata:`

`spec:`

---

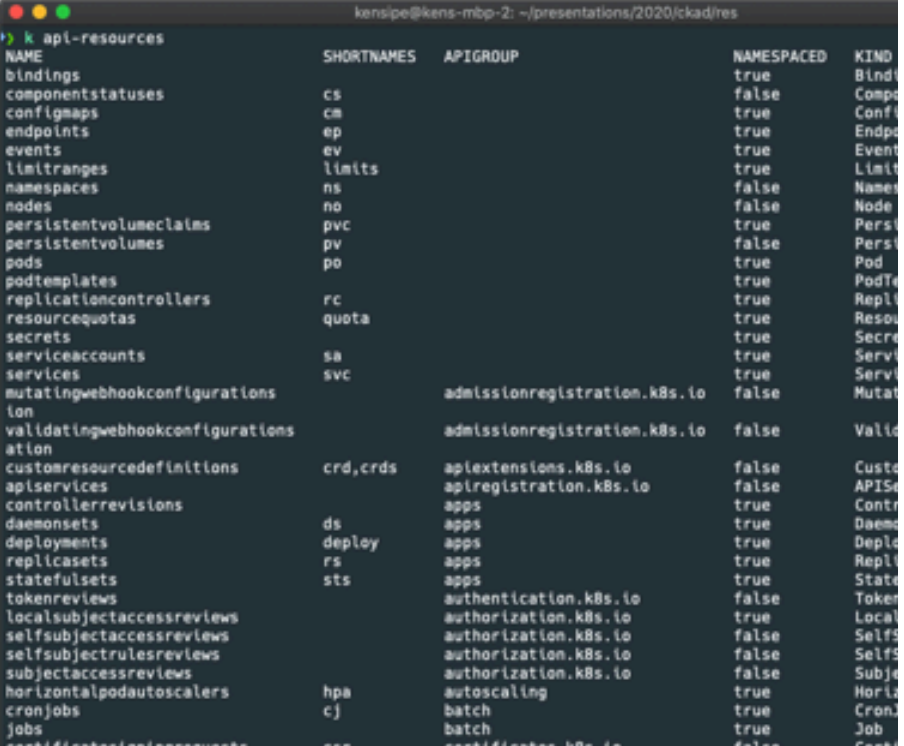
## Know Describe

- Describe pod, deployment, etc.
- Looking for annotations or events... reduce output

```
k describe pods nginx | grep --context=10 annotations:  
k describe pods | grep -C 10 Events:
```

# Know Shortcuts

- deploy - deployment
- svc - service
- ns - namespace
- netpol - networkPolicy
- pv, pvc, sa



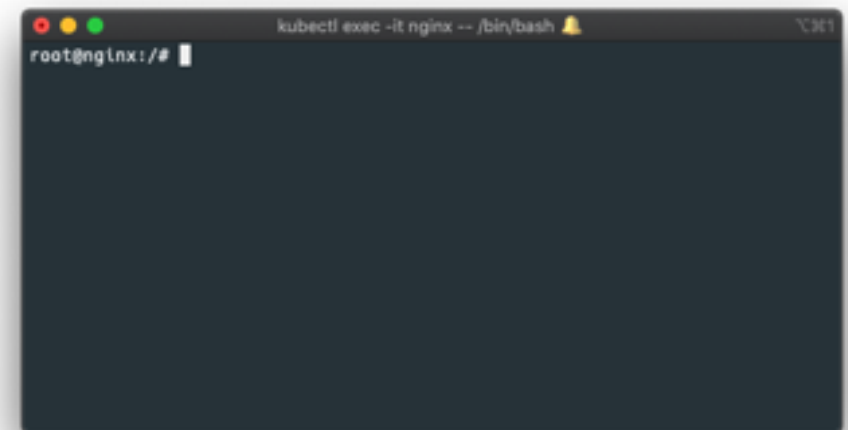
```
kensipe@kens-mbp-2: ~/presentations/2020/ckad/res
> kubectl api-resources
```

NAME	SHORTNAMES	APIGROUP	NAMESPACED	KIND
bindings			true	Bindi
componentstatuses	cs		false	Compo
configmaps	cm		true	Confi
endpoints	ep		true	Endpo
events	ev		true	Event
limitranges	limits		true	Limit
namespaces	ns		false	Names
nodes	no		false	Node
persistentvolumeclaims	pvc		true	Persi
persistentvolumes	pv		false	Persi
Pods	po		true	Pod
podtemplates			true	PodTe
replicationcontrollers	rc		true	Repli
resourcequotas	quota		true	Resou
secrets			true	Secre
serviceaccounts	sa		true	Servi
services	svc		true	Servi
mutatingwebhookconfigurations		admissionregistration.k8s.io	false	Mutat
validatingwebhookconfigurations		admissionregistration.k8s.io	false	Valid
customresourcedefinitions	crd,crds	apiregistration.k8s.io	false	Custo
apiservices		apiregistration.k8s.io	false	APIS
controllerrevisions		apps	true	Contr
daemonsets	ds	apps	true	Daemo
deployments	deploy	apps	true	Deplo
replicasets	rs	apps	true	Repli
statefulsets	sts	apps	true	State
tokenreviews		authentication.k8s.io	false	Token
localsubjectaccessreviews		authorization.k8s.io	true	Local
selfsubjectaccessreviews		authorization.k8s.io	false	Self
selfsubjectrulesreviews		authorization.k8s.io	false	Self
subjectaccessreviews		authorization.k8s.io	false	Subje
horizontalpodautoscalers	hpa	autoscaling	true	Horiz
cronjobs	cj	batch	true	Cron
jobs		batch	true	Job
certificatesigningrequests	csr	certificates.k8s.io	false	Certi

## Know how to enter a running Pod

- Assuming a pod named **nginx**

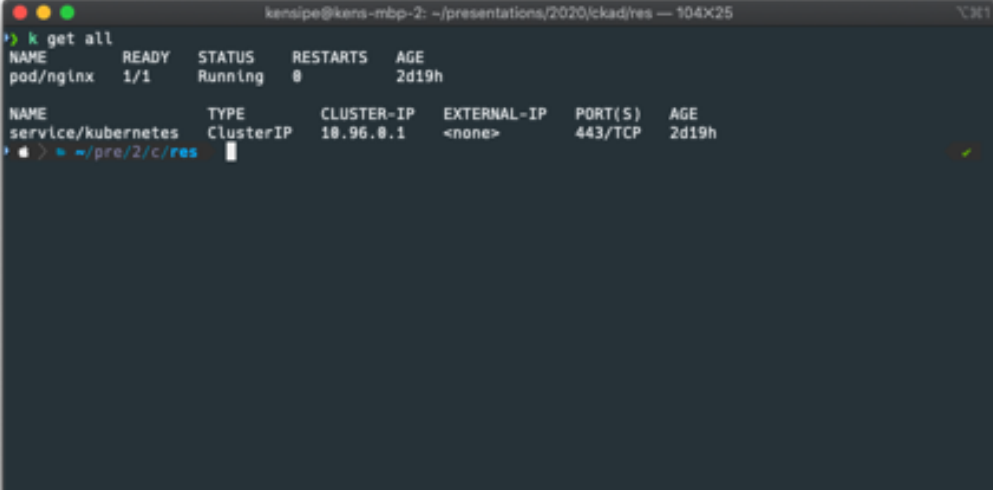
```
kubectl exec -it nginx -- /bin/bash
```





## Quick Namespace View

```
k get all
```

A terminal window titled 'kensipe@kens-mbp-2: ~/presentations/2020/ckad/res — 104x25'. The command 'k get all' has been executed, displaying two tables. The first table shows pod information: pod/nginx is in a 'Running' state with 1/1 ready containers, 0 restarts, and is 2d19h old. The second table shows service information: service/kubernetes is a 'ClusterIP' type with cluster IP 10.96.0.1, no external IP, port 443/TCP, and is 2d19h old. The prompt shows the user is in the directory ~/pre/2/c/res.

```
kensipe@kens-mbp-2: ~/presentations/2020/ckad/res — 104x25
> k get all
NAME          READY   STATUS    RESTARTS   AGE
pod/nginx     1/1     Running   0           2d19h

NAME          TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)    AGE
service/kubernetes  ClusterIP   10.96.0.1    <none>        443/TCP    2d19h
~ /pre/2/c/res
```

---

## Setting Context for Namespace

```
k config set-context $(k config current-context) --namespace=dev
```

```
k config set-context $(k config current-context) --namespace=default
```

## Quick View of Pod Env

```
k run busybox --image=busybox --command --restart=Never -- env
```

```
k logs busybox
PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
HOSTNAME=busybox
KUBERNETES_SERVICE_PORT_HTTPS=443
KUBERNETES_PORT=tcp://10.96.0.1:443
KUBERNETES_PORT_443_TCP=tcp://10.96.0.1:443
KUBERNETES_PORT_443_TCP_PROTO=tcp
KUBERNETES_PORT_443_TCP_PORT=443
KUBERNETES_PORT_443_TCP_ADDR=10.96.0.1
KUBERNETES_SERVICE_HOST=10.96.0.1
KUBERNETES_SERVICE_PORT=443
HOME=/root
```

---

## Edits

- `k edit (k edit pod nginx)`
- Prefer:
  - `K get pod nginx -o yaml > 1.yaml`
  - `vi 1.yaml`
  - `k apply -f 1.yaml`

---

## Speedy Deletes

- `k delete pod nginx --grace-period=1`
- `k delete pod nginx --grace-period=0 --force`

---

## Linux: Grep

- Grep
  - -i case-insensitive
  - -C # (content around match by # lines)

---

## Linux: One Liners

- Loops with output

```
while true; do date >> /var/log/app.txt; sleep 5; done
```

```
i=0; while true; do echo "$i: $(date)"; i=$((i+1));sleep 1; done
```

```
mkdir -p collect; while true; do cat /var/data/* > /collect/data.text; sleep 10; done
```

# Sample Apps

<https://github.com/kubernetes/examples>

README.md

## Kubernetes Examples

This directory contains a number of examples of how to run real applications with Kubernetes.

Refer to the [Kubernetes documentation](#) for how to execute the tutorials.

### Maintained Examples

Maintained Examples are expected to be updated with every Kubernetes release, to use the latest and greatest features, current guidelines and best practices, and to refresh command syntax, output, changed prerequisites, as needed.

Name	Description	Notable Features Used	Complexity Level
Guestbook	PHP app with Redis	Deployment, Service	Beginner
Guestbook-Go	Go app with Redis	Deployment, Service	Beginner
WordPress	WordPress with MySQL	Deployment, Persistent Volume with Claim	Beginner
Cassandra	Cloud Native Cassandra	Daemon Set, Stateful Set, Replication Controller	Intermediate

Note: Please add examples that are maintained to the list above.

See [Example Guidelines](#) for a description of what goes in this directory, and what examples should contain.



# Starting

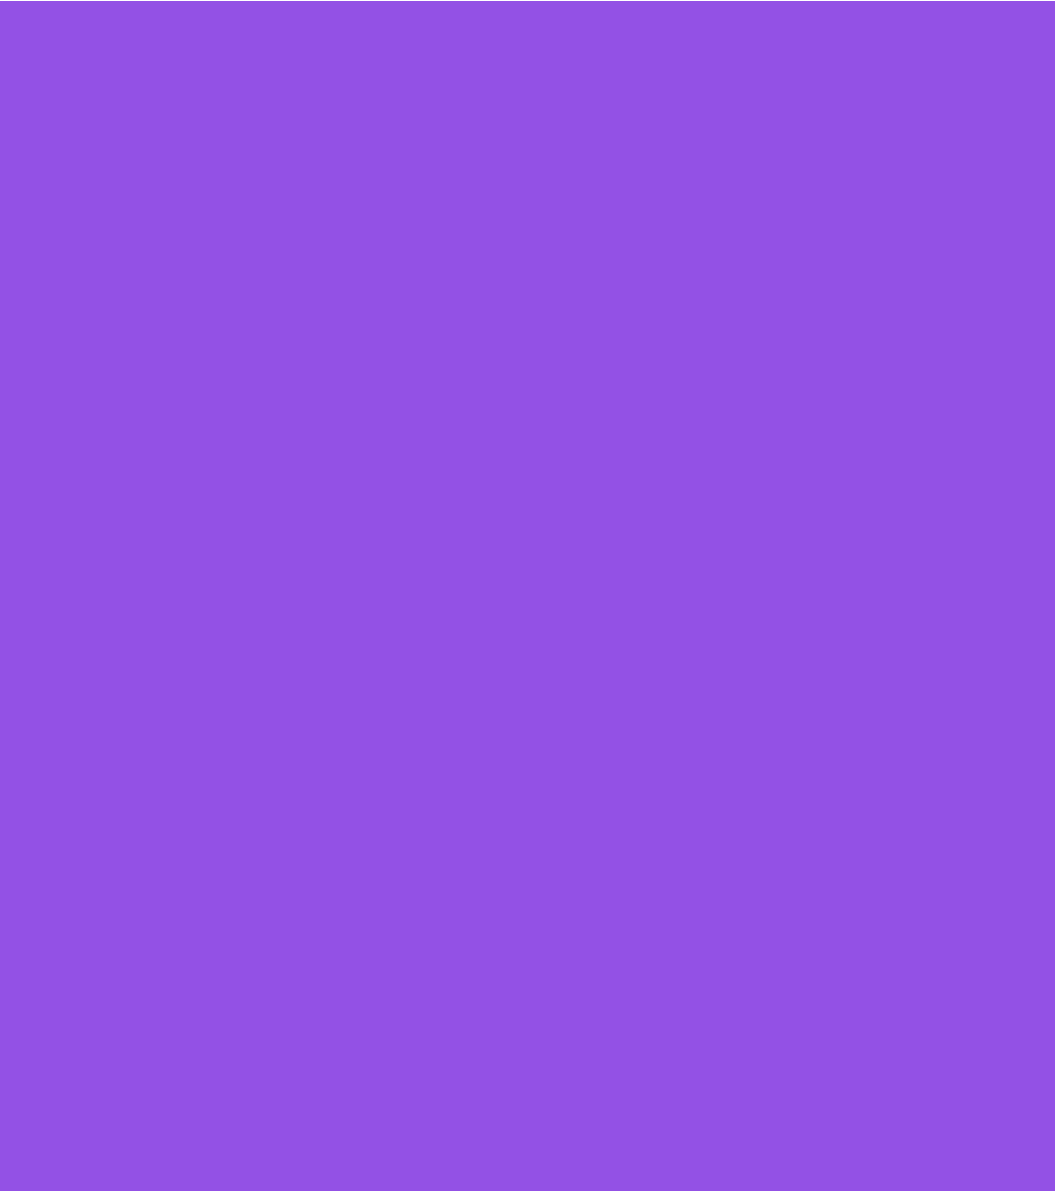
## Before Time Starts

- Clean Desktop
- 2 Browser (half screen each or 2 displays)
- Bookmarks for [k8s.io](https://k8s.io)

## Immediately when time starts

```
alias k=kubectl  
alias vi=vim
```

```
vi ~/.vimrc  
:set number  
:set et  
:set sw=2 ts=2 sts=2
```



---

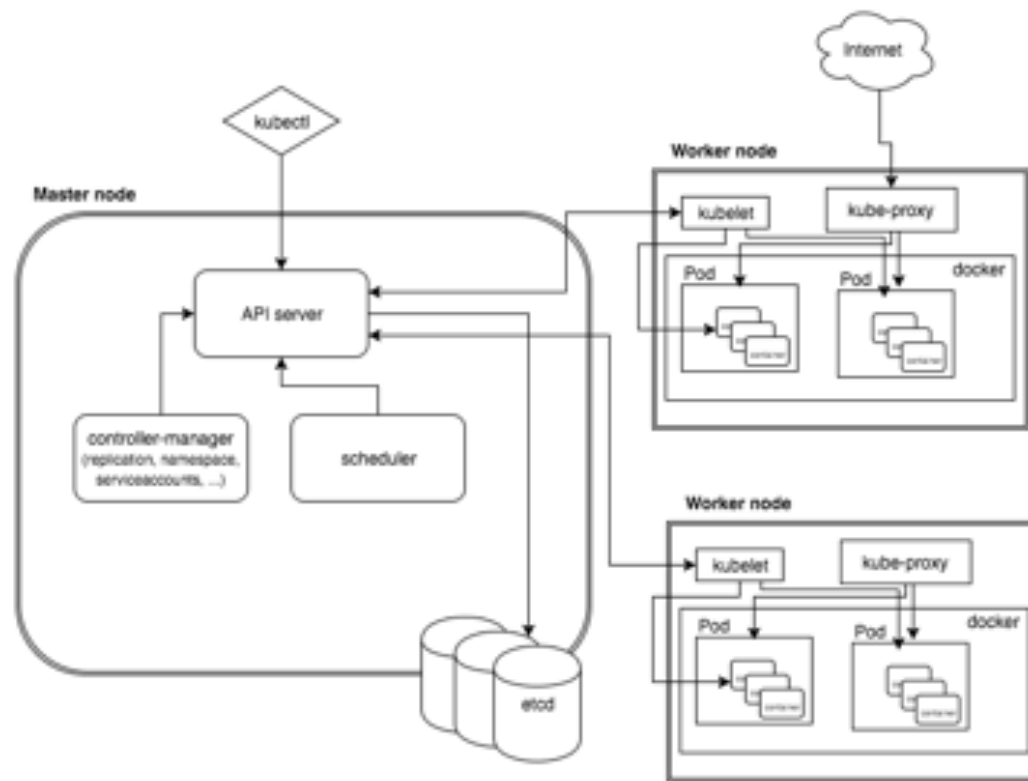
# Core Concepts 13%

---

# Core

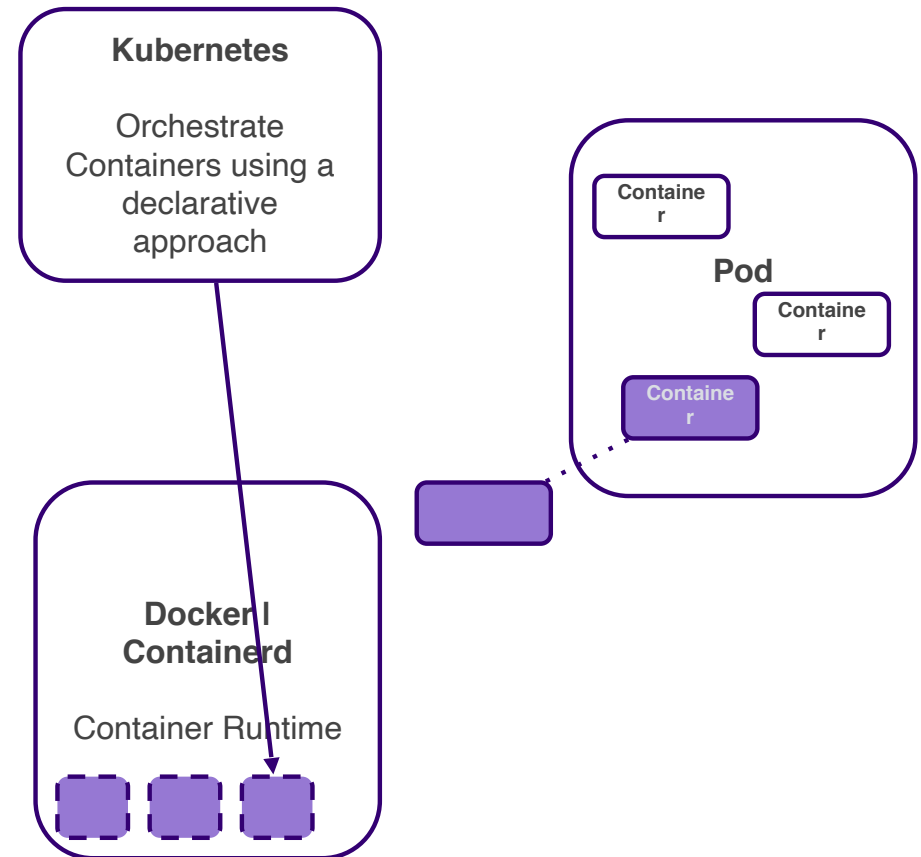
- Kube Architecture
- Pod
- YAML
- ReplicaSets
- Deployments
- Namespace

# KUBERNETES ARCHITECTURE



# Pods

- Pods provide an abstraction to
  - Automate scaling of containers
  - Adding resources such as networking and storage (shared resources)
  - Ability to use different container runtimes - we are using containerd
  - A Pod can contain one or more containers
- Kubernetes leverages Pods as the key abstraction to orchestrate containers
- All containers in a Pod live within a logical unit that has a single IP, but can communicate with each other over localhost.



---

# Pods

- Fundamental Unit of execution
- 1+ containers
- Metadata

```
k run nginx --image nginx -lapp=front-end
```

# Pods

➤ `k run nginx --image nginx -lapp=front-end --dry-run -o yaml`

apiVersion: v1

kind: Pod

metadata:

creationTimestamp: null

labels:

app: front-end

name: nginx

spec:

containers:

- image: nginx

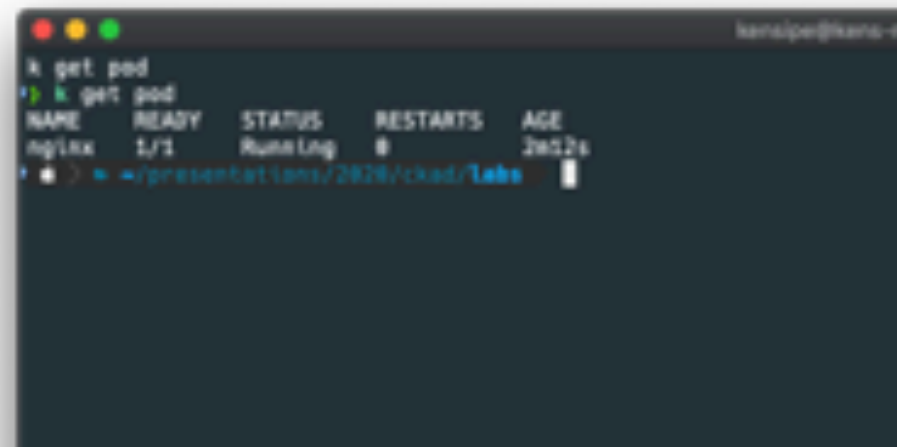
name: nginx

resources: {}

dnsPolicy: ClusterFirst

restartPolicy: Always

status: {}



```
k get pod
k get pod
NAME      READY   STATUS    RESTARTS   AGE
nginx     1/1     Running   0           2m12s
```

The screenshot shows a terminal window with a dark background. The prompt is 'k get pod' and the output is a table with columns: NAME, READY, STATUS, RESTARTS, and AGE. The table shows one pod named 'nginx' with 1/1 ready, Running status, 0 restarts, and an age of 2m12s. The terminal also shows the command 'k get pod' being entered again.

---

# ReplicaSet

- ReplicaSetController - Provides high availability
- Ensures specified number of pods are running across the cluster



# ReplicationController

```
k create -f 3app-3.yaml
replicationcontroller/webapp-rc created
k get rc
NAME          DESIRED   CURRENT   READY   AGE
webapp-rc     3         3         3       5s
k get pods
NAME          READY   STATUS    RESTARTS   AGE
nginx         1/1     Running   0          28h
webapp-rc-dwb4g 1/1     Running   0          9s
webapp-rc-kgxpd 1/1     Running   0          9s
```

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: webapp-rc
  labels:
    name: webapp
spec:
  template:
    metadata:
      name: webapp
      labels:
        app: front-end
    spec:
      containers:
        - name: nginx-container
          image: nginx
  replicas: 3
```

# ReplicaSet

```
k create -f lab9-1.yaml  
replicaset.apps/webapp-rs created
```

```
k get rs
```

NAME	DESIRED	CURRENT	READY	AGE
webapp-rs	3	3	3	6s

```
k get pods
```

NAME	READY	STATUS	RESTARTS	AGE
webapp-rs-6qdzk	1/1	Running	0	11s
webapp-rs-fcfbq	1/1	Running	0	11s
webapp-rs-grvg2	1/1	Running	0	11s

```
apiVersion: apps/v1  
kind: ReplicaSet  
metadata:  
  name: webapp-rs  
  labels:  
    name: webapp  
spec:  
  template:  
    metadata:  
      name: webapp  
    labels:  
      app: front-end  
  spec:  
    containers:  
      - name: nginx-container  
        image: nginx  
  replicas: 3  
  selector:  
    matchLabels:  
      app: front-end
```

## Scaling

- Update “replicas” in manifest yaml, then `k replace -f rs-def.yaml``
- `k scale --replicas= 6 -f rs-def.yaml`
- `k scale --replicas=6 rs webapp-rs`

---

## Kubectl commands

```
k create -f rs-def.yaml
```

```
k get rs
```

```
k delete -f rs-def.yaml
```

```
k replace -f rs-def.yaml
```

```
k scale --replicas=6 -f rs-def.yaml
```

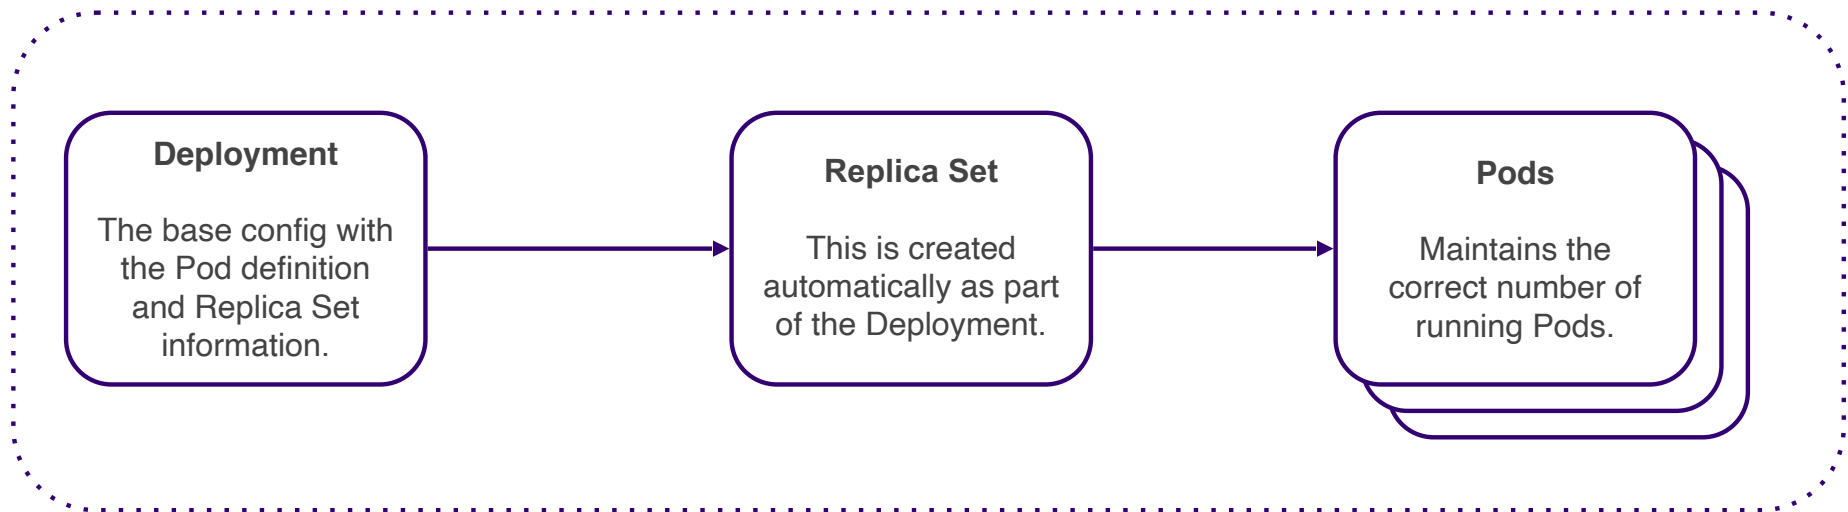
---

# Deployments

- A versioned set of replicaset, which is a collection of pods

# Deployments

- A versioned set of replicaset, which is a collection of pods
- They should be used instead of Replica Sets in most cases.
- Additional functionality is added in the form of rollouts and rollbacks.
- Allows developers and operations staff to scale up and scale down applications.
- Unlike a Replica Set it contains the Pod definition as part of its configuration.



## Deployment Manifest

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: webapp-rs
  labels:
    name: webapp
spec:
  template:
    metadata:
      name: webapp
      labels:
        app: front-end
    spec:
      containers:
        - name: nginx-container
          image: nginx
  replicas: 3
  selector:
    matchLabels:
      app: front-end
```



```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: webapp-deploy
  labels:
    name: webapp
spec:
  template:
    metadata:
      name: webapp
      labels:
        app: front-end
    spec:
      containers:
        - name: nginx-container
          image: nginx
  replicas: 3
  selector:
    matchLabels:
      app: front-end
```

# Deployments

```
k create -f 3app-2.yaml
deployment.apps/webapp-deploy created
k get deploy
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
webapp-deploy 3/3      3             3            6s
k get rs
NAME          DESIRED   CURRENT   READY   AGE
webapp-deploy-7f4cdb89 3          3          3        9s
k get pods
NAME          READY   STATUS    RESTARTS   AGE
webapp-deploy-7f4cdb89-52jvt 1/1     Running   0           14s
webapp-deploy-7f4cdb89-edrh9 1/1     Running   0           14s
webapp-deploy-7f4cdb89-p4b5z 1/1     Running   0           14s
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: webapp-deploy
  labels:
    name: webapp
spec:
  template:
    metadata:
      name: webapp
      labels:
        app: front-end
    spec:
      containers:
        - name: nginx-container
          image: nginx
      replicas: 3
      selector:
        matchLabels:
          app: front-end
```



## Deployments Status and History

```
k create -f lab9-2.yaml
deployment.apps/webapp-deploy created
k get deploy
NAME          READY   UP-TO-DATE   AVAILABLE   AGE
webapp-deploy 3/3      3            3           7s
k get rs
NAME          DESIRED   CURRENT   READY   AGE
webapp-deploy-7f4cdb89 3          3         3       14s
k get pods
NAME          READY   STATUS    RESTARTS   AGE
webapp-deploy-7f4cdb89-pf2zw 1/1     Running   0          18s
webapp-deploy-7f4cdb89-sqjgz 1/1     Running   0          18s
webapp-deploy-7f4cdb89-wtx9s 1/1     Running   0          18s
k rollout status deploy webapp-deploy
deployment "webapp-deploy" successfully rolled out
k rollout history deploy webapp-deploy
deployment.apps/webapp-deploy
REVISION   CHANGE-CAUSE
1          <none>
```

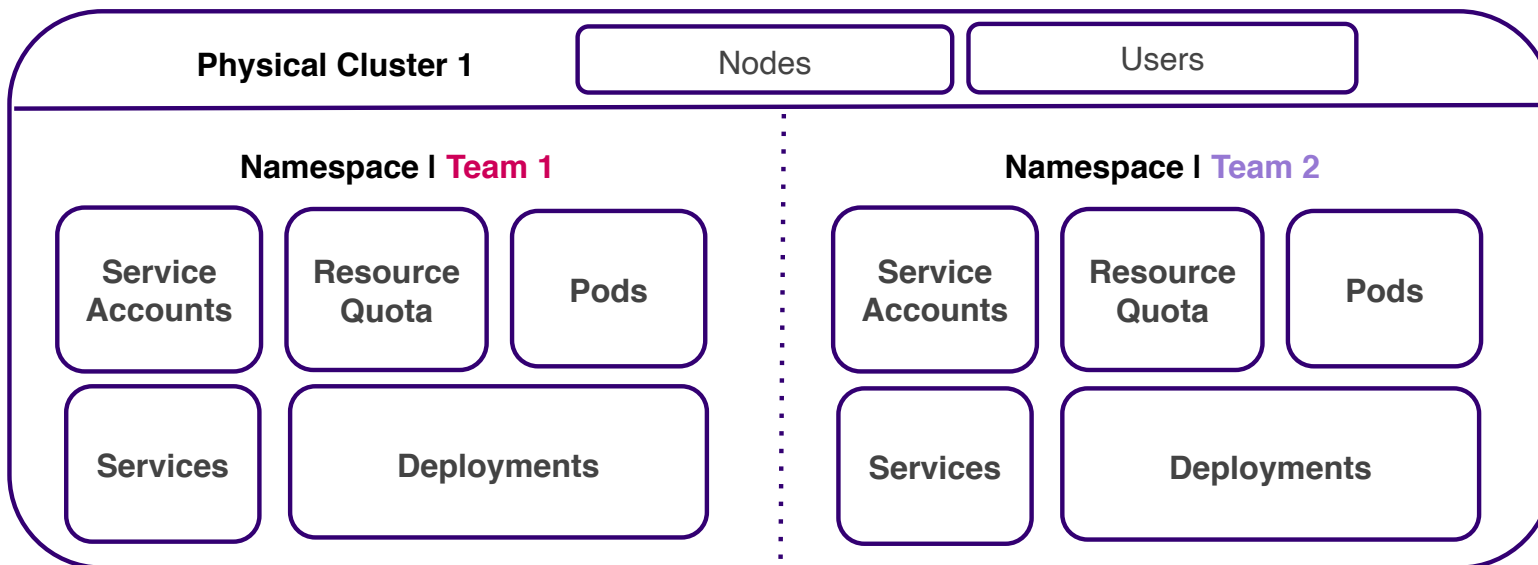
---

# Namespace

- Provides logical name separation
- default is the default ns
- kube-system is the system namespace

# Namespaces

- Each namespace can have its own set of policies which define:
  - Resource limits
  - Quota
  - Security



## Naming in Namespace and DNS

- Within a namespace, a service can “reference” a service by its name:

```
service.connect("db-service")
```

- References outside namespace, it is necessary to provide namespace details

```
service.connect("db-service.dev.svc.cluster.local")
```

 for a “db-service” in the “dev” namespace.

## Naming in Namespace and DNS

`db-service.dev.svc.cluster.local`

The diagram illustrates the components of the DNS name `db-service.dev.svc.cluster.local`. It features four horizontal blue lines positioned above the labels `db-service`, `dev`, `svc`, and `cluster.local`. Vertical blue lines connect these labels to their respective descriptions below: `db-service` is the service name, `dev` is the namespace, `svc` is the service, and `cluster.local` is the domain.

service name    namespace    service    domain

---

## Switching Namespaces

```
k config set-context $(k config current-context) --namespace=dev
```

Always know your context / namespace during the test!



---

# Configuration 18%

---

# Configuration

- Commands and Arguments
- ENV
- ConfigMaps
- Secrets
- Security Context
- Service Accounts
- Resource Limits
- Taints / Tolerations
- Affinity



## Commands

From Ubuntu

ENTRYPOINT ["sleep"]

CMD ["5"]



```
apiVersion: v1
kind: Pod
metadata:
  name: sleeper
spec:
  containers:
  - image: ubuntu-sleeper
    name: ubuntu
    args: ["10"]
```

docker run ubuntu-sleeper 10

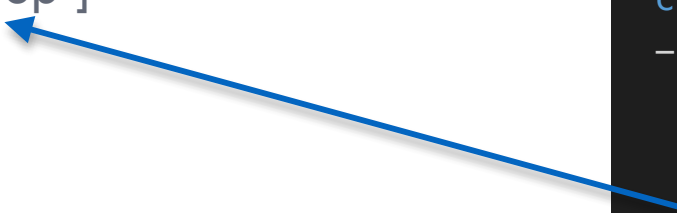
# Commands

From Ubuntu

ENTRYPOINT ["sleep"]

CMD ["5"]

```
apiVersion: v1
kind: Pod
metadata:
  name: sleeper
spec:
  containers:
  - image: ubuntu-sleeper
    name: ubuntu
    args: ["10"]
    command: ["sleep2.0"]
```



`docker run --entrypoint sleep2.0 ubuntu-sleeper 10`

# ENV

- Key / Value pairs
- Array under “env”

```
apiVersion: v1
kind: Pod
metadata:
  name: sleeper
spec:
  containers:
  - image: ubuntu-sleeper
    name: ubuntu
    env:
    - name: SLEEP_TIME
      value: 1200
```

# ConfigMaps

## Create ConfigMap

### •Imperative

```
k create configmap app-config --from-literal=SLEEP_TIME=1200
```

```
k create configmap app-config --from-file=app_config.properties
```

### •Declarative

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: app-config
data:
  SLEEP_TIME: "1200"
```

## Inject ConfigMap

```
apiVersion: v1
kind: Pod
metadata:
  name: ubuntu-sleeper-2
spec:
  containers:
    - image: ubuntu
      envFrom:
        - configMapRef:
            name: app-config
```

# Secrets

## Create Secret

### •Imperative

```
k create secret generic app-secret --from-literal=DB_Pass=mysecret
```

```
k create secret generic app-secret --from-file=app_secret.properties
```

### •Declarative

```
apiVersion: v1
kind: Secret
metadata:
  name: app-secret
data:
  DB_Password: bXlzZWNyZXQ=
```

linux: echo -n 'mysecret' | base64

## Inject Secret

```
apiVersion: v1
kind: Pod
metadata:
  name: ubuntu-sleeper-2
spec:
  containers:
    - image: ubuntu
      envFrom:
        - secretRef:
            name: app-secret
```

# Secrets

## View Secrets

```
k create secret generic app-secret --  
from-literal=DB_Pass=mysecret
```

```
> k describe secrets app-secret  
Name:          app-secret  
Namespace:     default  
Labels:        <none>  
Annotations:   <none>
```

```
Type:  Opaque
```

```
Data
```

```
====
```

```
DB_Pass:  8 bytes
```

view value: `k get secrets app-secret -o yaml`

linux: `echo -n 'bXlzZWNyZXQ=' | base64 --decode`

## File Based Secrets

- In container

ls /opt/app-secret-volume

DB\_Pass

```
spec:
  containers:
  - image: ubuntu
    name: ubuntu
    volumes:
    - name: app-secret-volume
      secret:
        secretName: app-secret
```

# Security Context

## Pod level security

```
apiVersion: v1
kind: Pod
metadata:
  name: ubuntu-sleeper
spec:
  securityContext:
    runAsUser: 1000
    capabilities:
      add: ["MAC_ADMIN"]
  containers:
    - image: ubuntu
      name: ubuntu
      command: ["sleep", "2000"]
```

## Container level security

```
apiVersion: v1
kind: Pod
metadata:
  name: ubuntu-sleeper
spec:
  containers:
    - image: ubuntu
      name: ubuntu
      command: ["sleep", "2000"]
      securityContext:
        runAsUser: 1000
        capabilities:
          add: ["MAC_ADMIN"]
```



---

## Security Accounts

- User

- Admin
- Developer

- Service

- Prometheus
- Jenkins

## Working with Service Accounts (sa)

- `k create sa foo`

- `k get sa`

NAME	SECRETS	AGE
default	1	2d23h
foo	1	6s

```
>k describe sa foo
```

```
Name:          foo
Namespace:     default
Labels:        <none>
Annotations:   <none>
Image pull secrets: <none>
Mountable secrets: foo-token-8prwq
Tokens:      foo-token-8prwq
Events:        <none>
```

sa creates token, token is stored in a secret

## Working with Service Accounts (sa)

```
k describe secret foo-token-8prwq
```

```
Name:      foo-token-8prwq
```

```
Namespace:  default
```

```
....
```

```
Data
```

```
====
```

```
ca.crt:    1025 bytes
```

```
namespace: 7 bytes
```

```
token:
```

```
eyJhbGciOiJSUzI1NiIsImtpZCI6ImlhbmVWVh5RldSTEEdwSkRtZWFWakdoU0hUZk5tNWxqTzhqcmRJKkQ1M0RidFEifQ.eyJpc3MiOiJrdWJlcm5ldGVzL3NlcnZpY2VhY2NvdW50liwia3ViZXJuZXRlcy5pby9zZXJ2aWNIYWVWb3VudC9uYW1lc3BhY2UiOiJkZWZhdWx0liwia3ViZXJuZXRlcy5pby9zZXJ2aWNIYWVWb3VudC9zZW5yZXQubmFtZSI6ImZvby10b2t1bi04cHJ3cSIsImt1YmVybmV0ZXMuaW8vc2VydmliZWFjY291bnQvc2VydmliZS1hY2NvdW50Lm5hbWUiOiJmb28iLCJrdWJlcm5ldGVzLmlvL3NlcnZpY2VhY2NvdW50L3NlcnZpY2UtYWVWb3VudC51aWQiOiJjOWE2M2MxMS04YzFiLTRIYTUtODY1My1hZjRkZTZmOWUzM2QiLCJzdWliOiJzeXN0ZW06c2VydmliZWFjY291bnQ6ZGVmYXVsdDpmb28ifQ.egg4hKgZoeYadrqpcZXZAsf3ZXMEQScydHzA4lez7939p2FOQ79vfB2D7RND_8ClabuB5CstY70YIfKMI4ktBVdEYgetHlb4g3FbAMeoSLqVKAnTJKGboMaTXoJ_BBIO-
```

# Pods, Namespaces, SA

> k describe pod ubuntu-sleeper

Name: ubuntu-sleeper  
Namespace: default  
Node: kind-control-plane/172.17.0.2  
Labels: <none>  
Annotations: Status: Running  
Containers:  
 ubuntu:  
 Image: ubuntu  
 Command:  
 sleep  
 4800  
 State: Running  
 Started: Mon, 17 Aug 2020 15:12:43 -0500  
 Ready: True  
 Restart Count: 0  
 Environment: <none>  
 Mounts:  
 /var/run/secrets/kubernetes.io/serviceaccount from default-token-jk9mk (ro)

```
k exec -it ubuntu-sleeper -- /bin/bash
root@ubuntu-sleeper:/# ll /var/run/secrets/kubernetes.io/serviceaccount/
..2020_08_17_20_12_37.030576906/ namespace
..data/ token
ca.crt
root@ubuntu-sleeper:/# ll /var/run/secrets/kubernetes.io/serviceaccount/
total 4
drwxrwxrwt 3 root root 148 Aug 17 20:12 ./
drwxr-xr-x 3 root root 4096 Aug 17 21:32 ../
drwxr-xr-x 2 root root 100 Aug 17 20:12 ..2020_08_17_20_12_37.030576906/
lrwxrwxrwx 1 root root 31 Aug 17 20:12 ..data -> ..2020_08_17_20_12_37.030576906/
lrwxrwxrwx 1 root root 13 Aug 17 20:12 ca.crt -> ..data/ca.crt
lrwxrwxrwx 1 root root 16 Aug 17 20:12 namespace -> ..data/namespace
lrwxrwxrwx 1 root root 12 Aug 17 20:12 token -> ..data/token
root@ubuntu-sleeper:/#
```

# Resource

- Default Requirements for a Pod
- CPU: 0.5
- MeM: 256Mi

CPU 0.1 is min  
(also known as 100m)

```
apiVersion: v1
kind: Pod
metadata:
  name: ubuntu-sleeper
spec:
  containers:
  - image: ubuntu
    name: ubuntu
    command: ["sleep", "2000"]
    resources:
      requests:
        memory: "1Gi"
        cpu: 1
```

# Resource Limits

- Limits on resources

```
limits:  
  memory: "2Gi"  
  cpu: 2
```

CPU is throttled  
Mem limits will OOM Killed

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: ubuntu-sleeper  
spec:  
  containers:  
  - image: ubuntu  
    name: ubuntu  
    command: ["sleep", "2000"]  
    resources:  
      requests:  
        memory: "1Gi"  
        cpu: 1  
      limits:  
        memory: "2Gi"  
        cpu: 2
```

# Node Taints

- Nodes can be tainted
- When a node is tainted, only tolerant pods can be scheduled.
- Anti-Affinity

```
k taint nodes node-name key=value:taint-effect
```

```
k taint nodes node1 app=red:NoSchedule
```

- Taint-Effect
  - NoSchedule
  - PreferNoSchedule
  - NoExecute

# Pod Tolerations

- Pods can have tolerations
- Used to tolerate a taint

```
kubectl describe node kubemaster | grep Taint  
Taints:          node-role.kubernetes.io/master:NoSchedule
```

```
apiVersion: v1  
kind: Pod  
metadata:  
  name: ubuntu-sleeper  
spec:  
  containers:  
  - image: ubuntu  
    name: ubuntu  
    command: ["sleep", "2000"]  
  tolerations:  
  - key: "app"  
    operator: "Equal"  
    value: "red"  
    effect: "NoSchedule"
```



# Node Selectors

- Label Nodes

```
k label nodes <node-name> key=value
```

```
k label nodes node-1 size=Large
```

- Limited:

- NOT Small
- Medium or Large

```
apiVersion: v1
kind: Pod
metadata:
  name: ubuntu-sleeper
spec:
  containers:
    - image: ubuntu
      name: ubuntu
      command: ["sleep", "2000"]
  nodeSelector:
    size: Large
```

# Node Affinity

## ●Affinity

- requiredDuringSchedulingIgnoredDuringExecution
- preferredDuringSchedulingIgnoredDuringExecution

## ●Operators

- In
- NotIn
- Exists
- DoesNotExist
- Gt
- Lt

```
apiVersion: v1
kind: Pod
metadata:
  name: ubuntu-sleeper
spec:
  containers:
    - image: ubuntu
      name: ubuntu
      command: ["sleep", "2000"]
  affinity:
    nodeAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        nodeSelectorTerms:
          - matchExpressions:
              - key: size
                operator: In
                values:
                  - Large
```



---

# Pod Design 20%

---

## Pods Design

- Labels, Selectors and Annotations
- Deployments
  - Rolling Updates
  - Rollbacks
- Jobs and CronJobs

## Labels and Selectors

- Labels provide a way to describe an object in a way that it can be selected

- app=front-end
- app=auth
- app=db

```
apiVersion: v1
kind: Pod
metadata:
  name: ubuntu-sleeper
  labels:
    app: front-end
spec:
  containers:
    - image: ubuntu
      name: ubuntu
      command: ["sleep", "2000"]
```

## Labels and Selectors

- Learn to create with generator

```
> k run nginx --image nginx -lapp=front-end
```

- Use Selector

```
> k get pods --selector app=front-end
```

```
apiVersion: v1
kind: Pod
metadata:
  name: ubuntu-sleeper
  labels:
    app: front-end
spec:
  containers:
    - image: ubuntu
      name: ubuntu
      command: ["sleep", "2000"]
```

# ReplicaSets

- ReplicaSet selector will match existing matching apps that
- Labels in template is for pod
- Labels for replicaset are separate
- Similar for service

```
apiVersion: v1
kind: ReplicaSet
metadata:
  name: ubuntu-sleeper
  labels:
    app: front-end
spec:
  replicas: 3
  selector:
    matchLabels:
      app: front-end
  template:
    metadata:
      labels:
        app: front-end
    spec:
      containers:
        - name: ubuntu-sleeper
          image:
```

# Annotations

- Annotations are not selectable

```
apiVersion: v1
kind: ReplicaSet
metadata:
  name: ubuntu-sleeper
  labels:
    app: front-end
  annotations:
    build: 1.0
spec:
  replicas: 3
  selector:
    matchLabels:
      app: front-end
  template:
```



# Deployments

- A new deployment, triggers a rollout
- Rollout creates a new revision
- Commands to know:
  - k rollout status deployment foo
  - k rollout history deployment foo
- Deployment Strategies
  - Recreate
  - Rolling Update (default)

# Deployments

```
k rollout history deploy web-app
deployment.apps/web-app
REVISION  CHANGE-CAUSE
1          <none>
```

```
k set image deploy web-app nginx-container=nginx:1.9.1
deployment.apps/web-app image updated
```

Events:				
Type	Reason	Age	From	Message
Normal	ScalingReplicaSet	5m7s	deployment-controller	Scaled up replica set web-app-597c9dfcfb to 3
Normal	ScalingReplicaSet	36s	deployment-controller	Scaled up replica set web-app-5bf48794c7 to 1
Normal	ScalingReplicaSet	17s	deployment-controller	Scaled down replica set web-app-597c9dfcfb to 2
Normal	ScalingReplicaSet	17s	deployment-controller	Scaled up replica set web-app-5bf48794c7 to 2

```
k describe deploy web-app
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: web-app
  labels:
    app: front-end
spec:
  replicas: 3
  selector:
    matchLabels:
      app: front-end
  template:
    metadata:
      name: web-app-pod
      labels:
        app: front-end
    spec:
      containers:
        - name: nginx-container
          image: nginx
```

All Rights Reserved.

# Deployments Imperatively

```
> k create deployment nginx --image=nginx
```

```
> k set image deploy nginx nginx=nginx:1.9.1  
deployment.apps/nginx image updated
```

```
> k describe deploy nginx  
Name: nginx  
Namespace: default  
CreationTimestamp: Tue, 18 Aug 2020 16:34:56 -0500  
Labels: app=nginx  
Annotations: deployment.kubernetes.io/revision: 2  
Selector: app=nginx  
Replicas: 1 desired | 1 updated | 1 total | 1 available | 0 unavailable  
StrategyType: RollingUpdate  
MinReadySeconds: 0  
RollingUpdateStrategy: 25% max unavailable, 25% max surge  
Pod Template:  
  Labels: app=nginx  
  Containers:  
    nginx:  
      Image: nginx:1.9.1  
      Port: <none>  
      Host Port: <none>  
      Environment: <none>  
      Mounts: <none>  
      Volumes: <none>  
Conditions:  
  Type      Status Reason  
  ----      -  
  Available True   MinimumReplicasAvailable  
  Progressing True  NewReplicaSetAvailable  
OldReplicaSets: <none>  
NewReplicaSet: nginx-7864b4477b (1/1 replicas created)  
Events:  
  Type      Reason      Age    From          Message  
  ----      -  
  Normal    ScalingReplicaSet  3m12s  deployment-controller  Scaled up replica set nginx-86c57db685 to 1  
  Normal    ScalingReplicaSet  112s   deployment-controller  Scaled up replica set nginx-7864b4477b to 1  
  Normal    ScalingReplicaSet  111s   deployment-controller  Scaled down replica set nginx-86c57db685 to 0
```

# Rollback

- Undo a deployment

```
k rollout undo deploy web-app
```

```
> k get rs
NAME                                DESIRED  CURRENT  READY  AGE
web-app-597c9dfcfb                 3         3        3     9m46s
web-app-5bf48794c7                 0         0        0     5m15s
```

# Jobs

```
kensipe@kens-mbp:~$ k create -f lab5-1.yaml
job.batch/addition-job created
kensipe@kens-mbp:~$ k get jobs
NAME          COMPLETIONS  DURATION  AGE
addition-job  1/1          9s        33s
kensipe@kens-mbp:~$ k get pod
NAME          READY  STATUS   RESTARTS  AGE
addition-job-555zh  0/1    Completed  0          48s
kensipe@kens-mbp:~$ k logs pod addition-job-555zh
Error from server (NotFound): pods "pod" not found
kensipe@kens-mbp:~$ k logs addition-job-555zh
5
```

```
apiVersion: batch/v1
kind: Job
metadata:
  name: addition-job
spec:
  template:
    spec:
      containers:
      - name: addition
        image: ubuntu
        command: ["expr", "3", "+", "2"]
        restartPolicy: OnFailure
```

# Jobs

- Completions - The number of successful completions before job is done
- Parallelism - The number of concurrent jobs at any given point in time.

```
apiVersion: batch/v1
kind: Job
metadata:
  name: addition-job
spec:
  completions: 3
  parallelism: 2
  template:
    spec:
      containers:
        - name: addition
          image: ubuntu
          command: ["expr", "3", "+", "2"]
          restartPolicy: OnFailure
```

# CronJob

- Spec inside a template, inside a spec, inside a jobTemplate, inside a spec!

- Schedule == cron

<https://en.wikipedia.org/wiki/Cron>

```
apiVersion: batch/v1beta1
kind: CronJob
metadata:
  name: addition-cron
spec:
  schedule: "*/1 * * * *"
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: addition
              image: ubuntu
              command: ["expr", "3", "+", "2"]
          restartPolicy: Never
```

```
k create cronjob nginx --image=nginx --schedule="* * * * *" --dry-run -o yaml
```

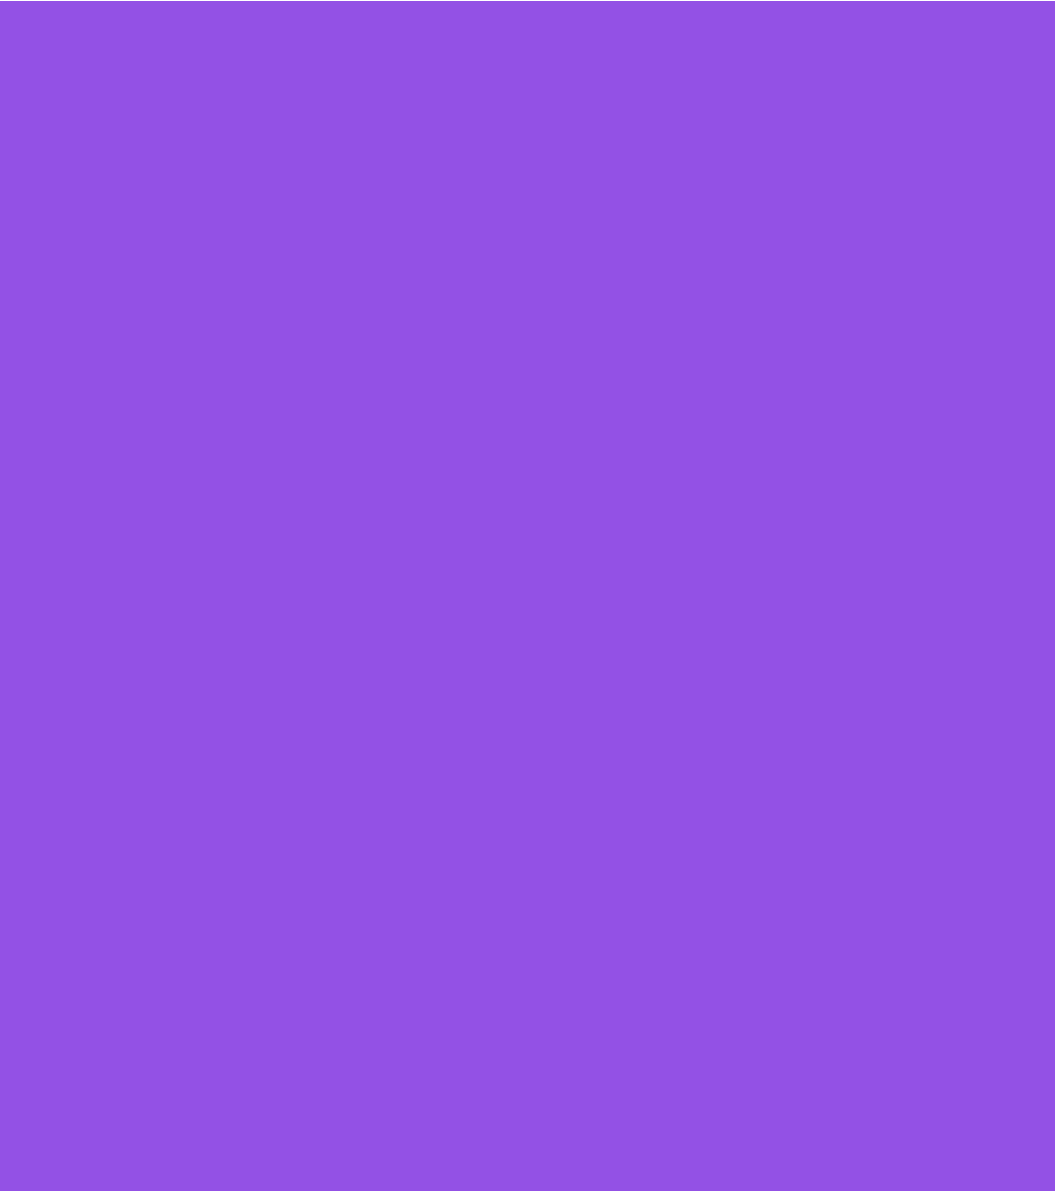
Rights Reserved.

## CronJob on Test

```
k create cronjob nginx --image=nginx --schedule="* * * * *" --dry-run -o yaml
```

- You can't google
- Know a few cron examples at least 5 positions





---

Observability  
18%

---

## Observability

- Container Logs
- LivenessProbes
- ReadinessProbes

# Container Logs

```
kensipe@kano-mbp-2: ~/presentations/2020/ckad/lab1
kubectl run box --image=busybox -- /bin/sh -c 'i=0; while true; do echo "i: $(date)"; i=$((i+1)); sleep 1; done'
> kubectl run box --image=busybox -- /bin/sh -c 'i=0; while true; do echo "i: $(date)"; i=$((i+1)); sleep 1; done'
pod/box created
> k get pod
NAME READY STATUS RESTARTS AGE
box 1/1 Running 0 3s
> k logs box
0: Wed Aug 19 20:11:30 UTC 2020
1: Wed Aug 19 20:11:31 UTC 2020
2: Wed Aug 19 20:11:32 UTC 2020
3: Wed Aug 19 20:11:33 UTC 2020
4: Wed Aug 19 20:11:34 UTC 2020
5: Wed Aug 19 20:11:35 UTC 2020
6: Wed Aug 19 20:11:36 UTC 2020
7: Wed Aug 19 20:11:37 UTC 2020
> k logs box
0: Wed Aug 19 20:11:38 UTC 2020
1: Wed Aug 19 20:11:39 UTC 2020
2: Wed Aug 19 20:11:40 UTC 2020
3: Wed Aug 19 20:11:41 UTC 2020
4: Wed Aug 19 20:11:42 UTC 2020
5: Wed Aug 19 20:11:43 UTC 2020
6: Wed Aug 19 20:11:44 UTC 2020
7: Wed Aug 19 20:11:45 UTC 2020
8: Wed Aug 19 20:11:46 UTC 2020
9: Wed Aug 19 20:11:47 UTC 2020
> k < ~/presentations/2020/ckad/lab1
```


## Pod Status / Conditions

- Pod Status
  - Pending - Trying to be scheduled
  - ContainerCreating - Pull Images
  - Run
- Pod Conditions (T/F)
  - PodScheduled
  - Initialized
  - ContainerReady
  - Ready

## Pod Status / Conditions

Condition Ready

Status



```
> k get pods
```

NAME	READY	STATUS	RESTARTS	AGE
box	1/1	Running	0	19m

---

## Meaning of Ready

- When a Container is “Ready” - Service Traffic will be routed to it!

---

## Readiness Probes

- HTTP
- TCP
- Exec Command (exit 0)

## HTTP Readiness Probe

```
apiVersion: v1
kind: Pod
metadata:
  name: webapp
  labels:
    name: webapp
spec:
  containers:
  - name: webapp
    image: nginx
    ports:
      - containerPort: 80
    readinessProbe:
      httpGet:
        path: /
        port: 80
```



# Readiness Probes

```
readinessProbe:  
  httpGet:  
    path: /  
    port: 80
```

```
readinessProbe:  
  tcpSocket:  
    port: 3306
```

```
readinessProbe:  
  exec:  
    command:  
      - cat  
      - /var/log/app.log
```

# Readiness Probes

```
readinessProbe:  
  httpGet:  
    path: /  
    port: 80  
  initialDelaySeconds: 10  
  periodSeconds: 5  
  failureThreshold: 8
```

```
readinessProbe:  
  tcpSocket:  
    port: 3306
```

```
readinessProbe:  
  exec:  
    command:  
      - cat  
      - /var/log/app.log
```

---

# Liveness Probes

Health Check!

- HTTP
- TCP
- Exec Command (exit 0)

# HTTP Liveness Probe

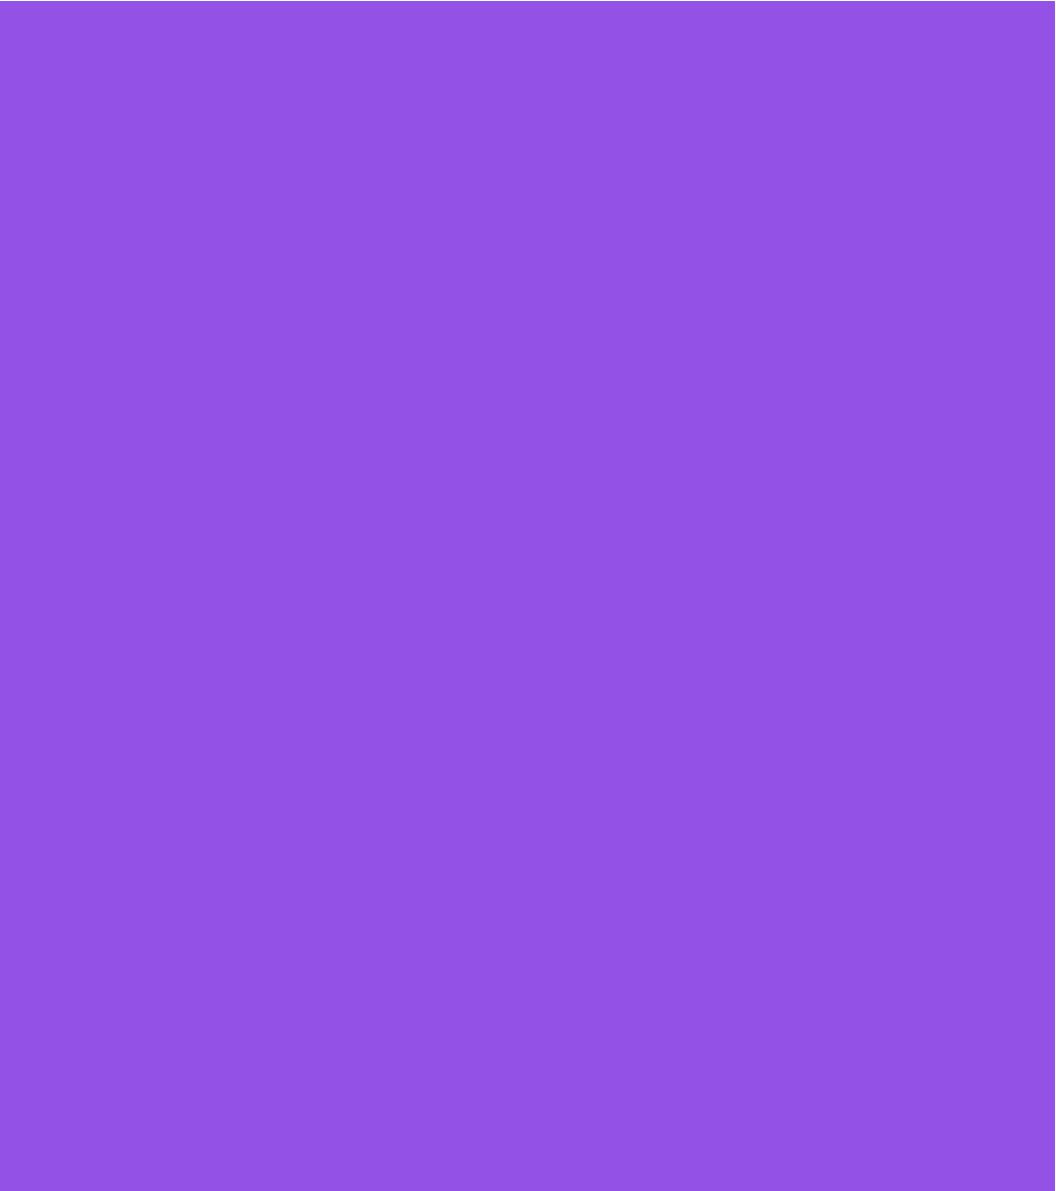
```
apiVersion: v1
kind: Pod
metadata:
  name: webapp
  labels:
    name: webapp
spec:
  containers:
  - name: webapp
    image: nginx
    ports:
      - containerPort: 80
    livenessProbe:
      httpGet:
        path: /
        port: 80
```

# Liveness Probes

```
livenessProbe:  
  httpGet:  
    path: /  
    port: 80
```

```
livenessProbe:  
  tcpSocket:  
    port: 3306
```

```
livenessProbe:  
  exec:  
    command:  
      - cat  
      - /var/log/app.log
```



---

# Multi- Container 10%

---

## Multi-Container

- Ambassador
- Adapter
- Sidecar

---

# Pods

- Designed to be Multi-Container
- Share:
  - Lifecycle
  - Network
  - Storage



---

## Multi-Container

```
apiVersion: v1
kind: Pod
metadata:
  name: webapp
  labels:
    name: webapp
spec:
  containers:
  - name: webapp
    image: nginx
    ports:
      - containerPort: 80
  - name: log-agent
    image: log-agent
```

---

# Architectural Patterns

- **Sidecar**

- 2nd non-primary container for logging, analytics, etc.

- **Adapter**

- 2nd non-primary container that processes to common format or summary format which forwards to another service

- **Ambassador**

- 2nd container that switches env your are connecting to... like dev, test, prod



---

Services /  
Networking  
13%

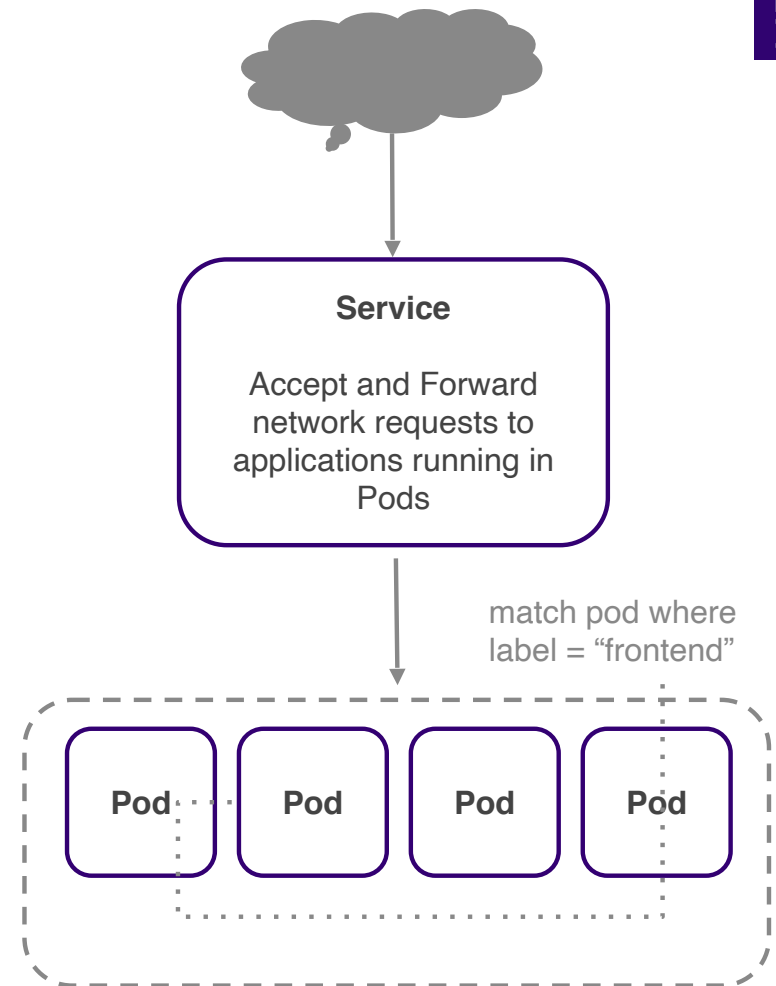
---

## Services / Networking

- Services
- NetworkPolicies

# Services

- Networking rules based on labels to control traffic to applications
- Service type determines how the application is exposed
  - Within the cluster (**ClusterIP**)
  - Through a static port on each node which allows external access (**NodePort**)
  - Through load balancer - physical or cloud provider (**LoadBalancer**)



---

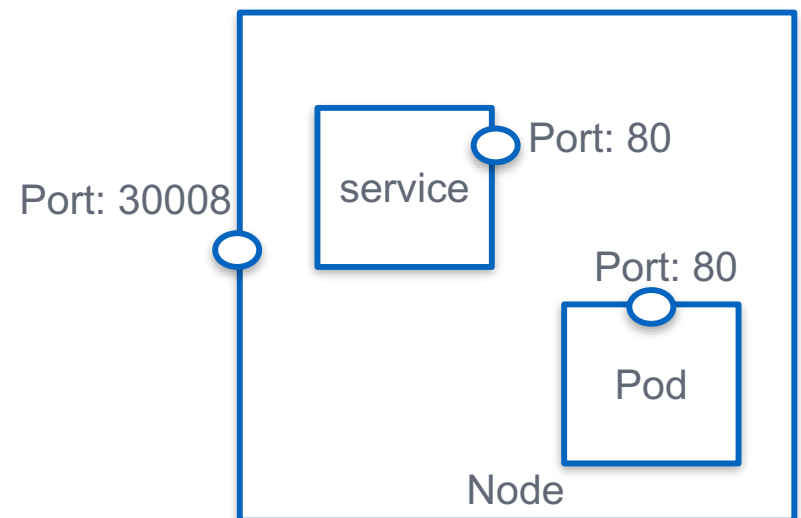
## Service Types

- NodePort - service an internal pod port available on the node
- ClusterIP - creates a VIP in the cluster to enable access to a set of pods for a service (think backend or mid-tier service like a database)
- LoadBalancer - provisions a LB from cloud providers (think front-end service like web)

## Service - NodePort

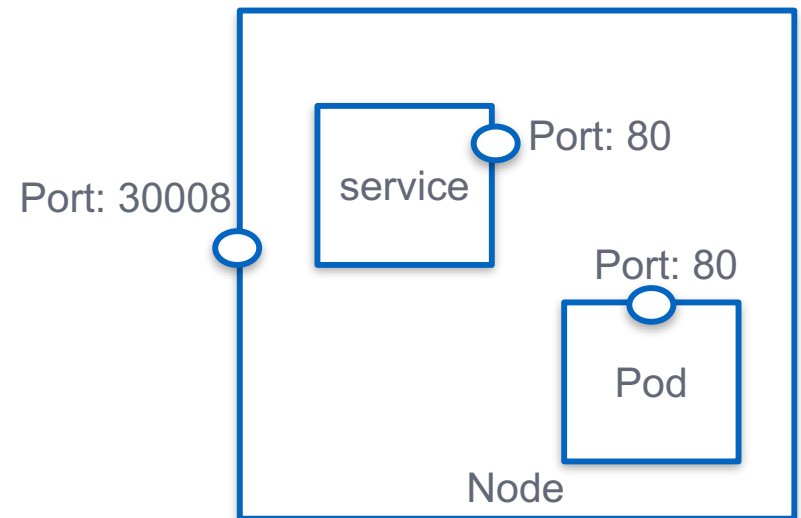
- Target Port - on pod: 80
- Port - on service
- NodePort - on node: 30008
  - default range: 30,000 - 32,767

\* terms are from the perspective of the service



## Service - NodePort

```
apiVersion: v1
kind: Service
metadata:
  name: app-service
spec:
  type: NodePort
  ports:
    - targetPort: 80
      port: 80
      nodePort: 30008
  selector:
    app: front-end
```

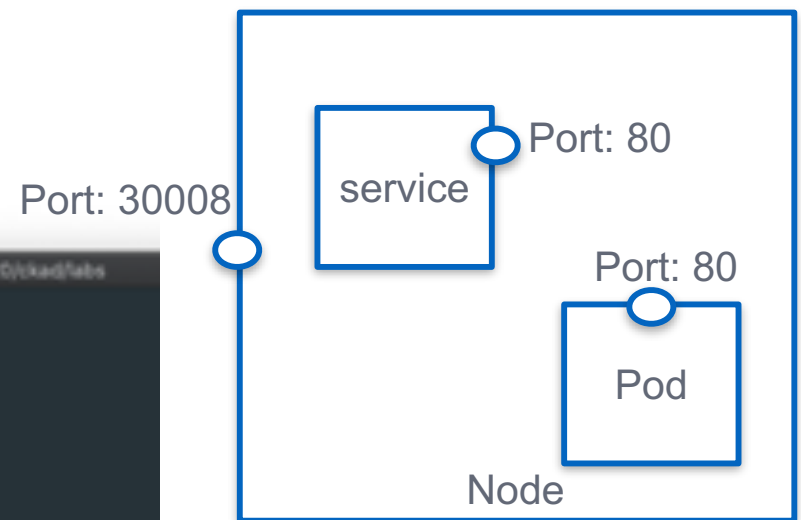


- If you don't provide a targetPort, it is assumed to be the same as port
- If you don't provide a nodePort, it will be provided within the range



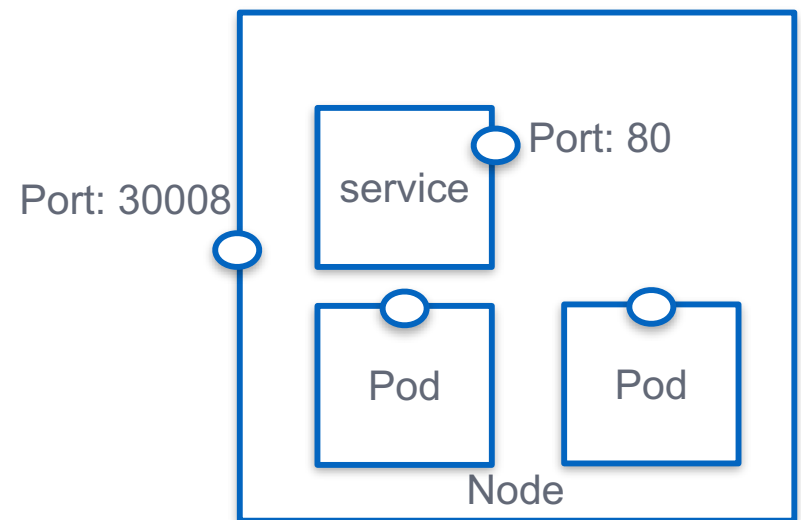
## Service - NodePort

```
kensipe@kens-mbp-2: ~/presentations/2020/ckad/labs
> k run nginx --image nginx --lapp=front-end
pod/nginx created
> k create -f svc-def.yaml
service/app-service created
> k get svc
NAME          TYPE          CLUSTER-IP    EXTERNAL-IP    PORT(S)          AGE
app-service   NodePort      10.96.225.237 <none>         80:30008/TCP     4s
kubernetes    ClusterIP     10.96.0.1     <none>         443/TCP          16h
```



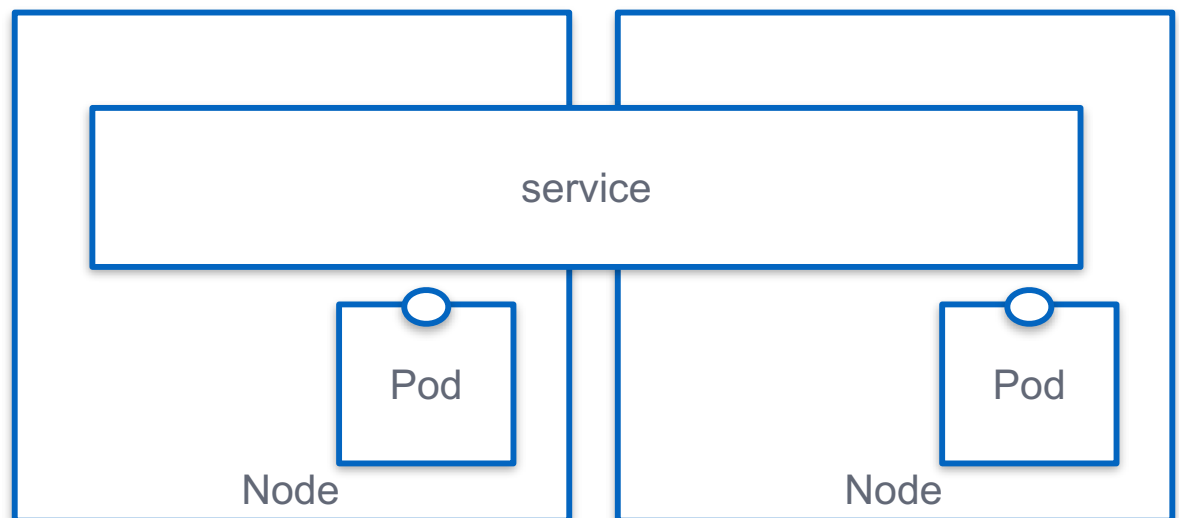
## Service - NodePort

- Multiple pods on Node are:
- Randomly load balanced
- With Session Affinity



## Service - NodePort

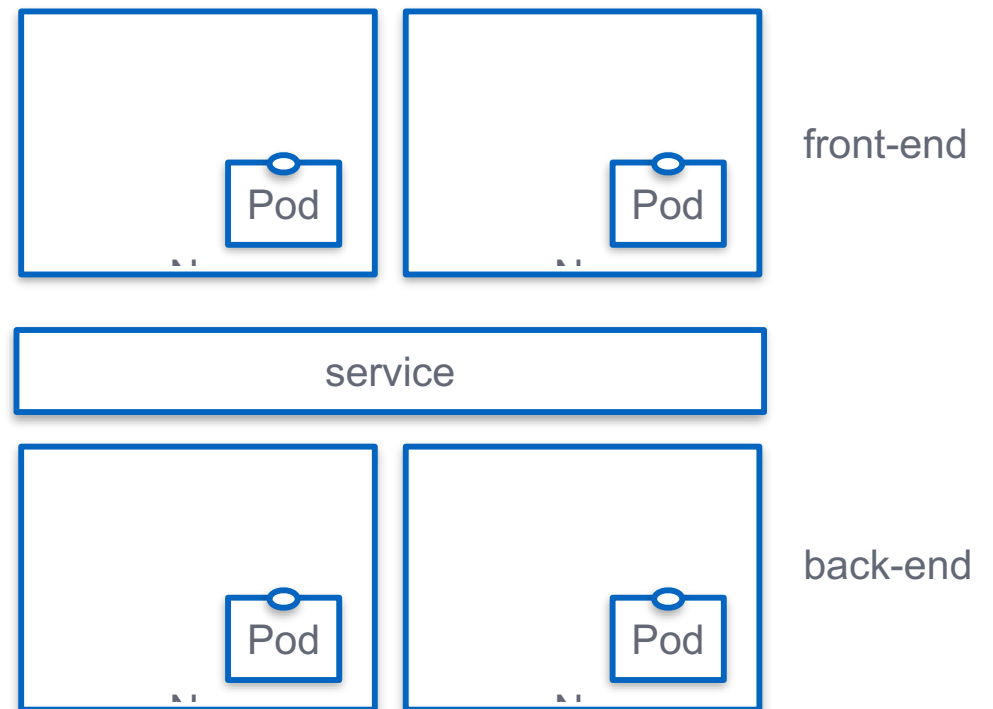
- Logic Service across Nodes
- All Node IPs have same NodePort port to provide access to service



## Service - ClusterIP

- Pods have IP... but they are elastic

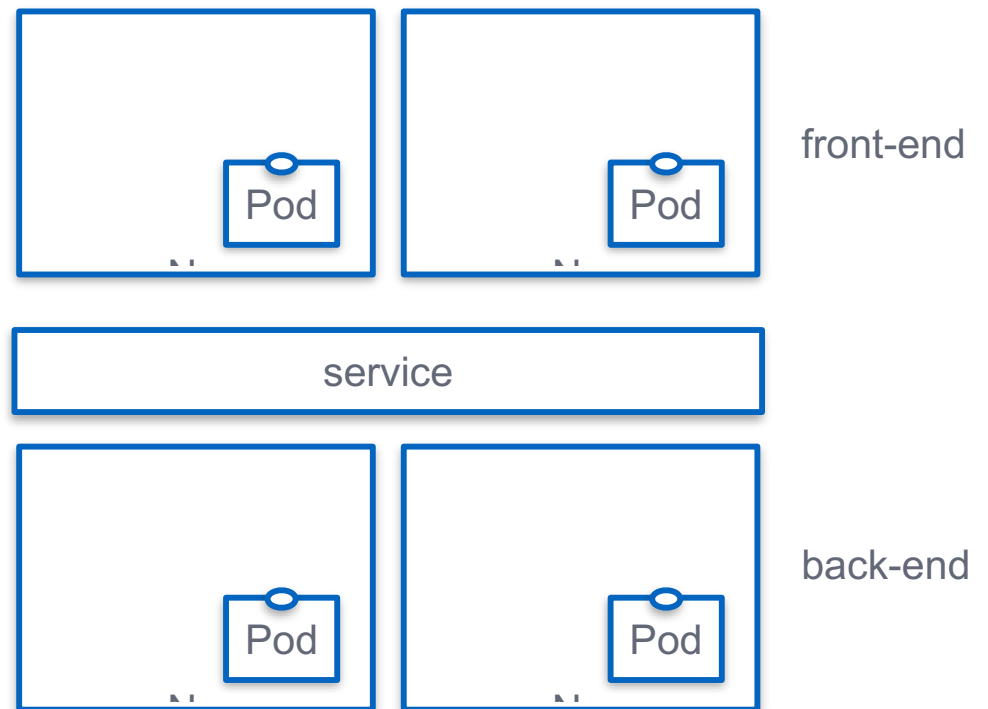
- 



## Service - ClusterIP

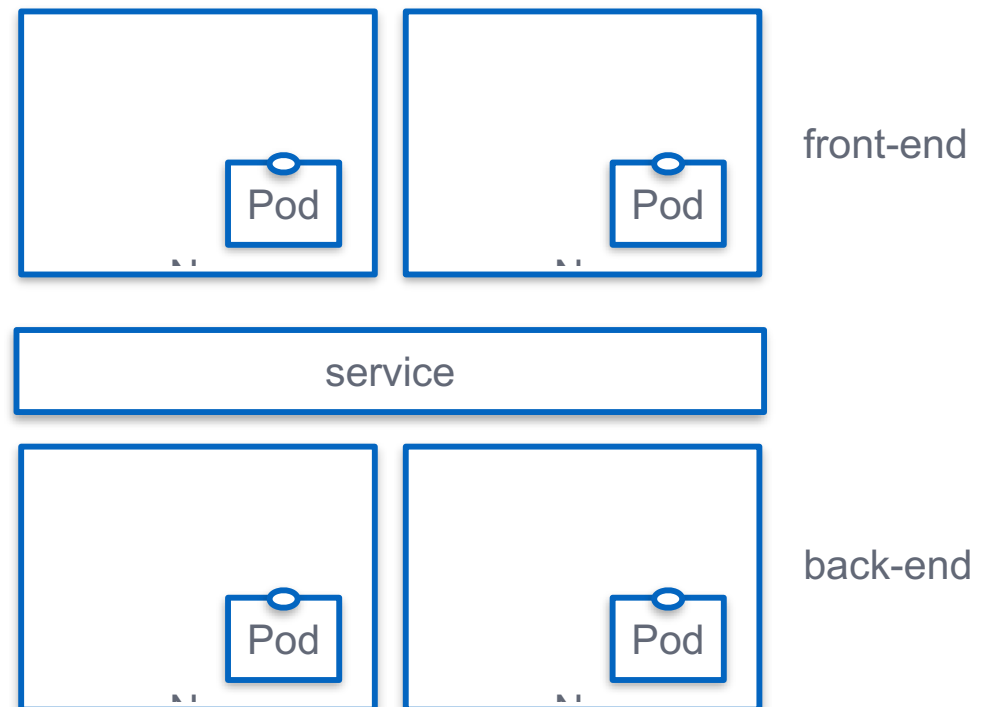
```
apiVersion: v1
kind: Service
metadata:
  name: back-end
spec:
  type: ClusterIP
  ports:
    - targetPort: 80
      port: 80
  selector:
    app: back-end
```

- type: ClusterIP is the default



## Service - ClusterIP

```
kensipe@kens-mbp-2: ~/presentations/2020/ckad/labs
> k run nginx --image nginx --lapp=back-end
pod/nginx created
> k create -f svc-def.yaml
service/back-end created
> k get svc
NAME      TYPE      CLUSTER-IP   EXTERNAL-IP   PORT(S)   AGE
back-end  ClusterIP  10.96.42.25   <none>        80/TCP    3s
kubernetes ClusterIP  10.96.0.1     <none>        443/TCP   103s
```



---

## Service - ClusterIP / NodePort

- ClusterIP

k expose pod nginx --port

k expose deploy nginx --port

- NodePort

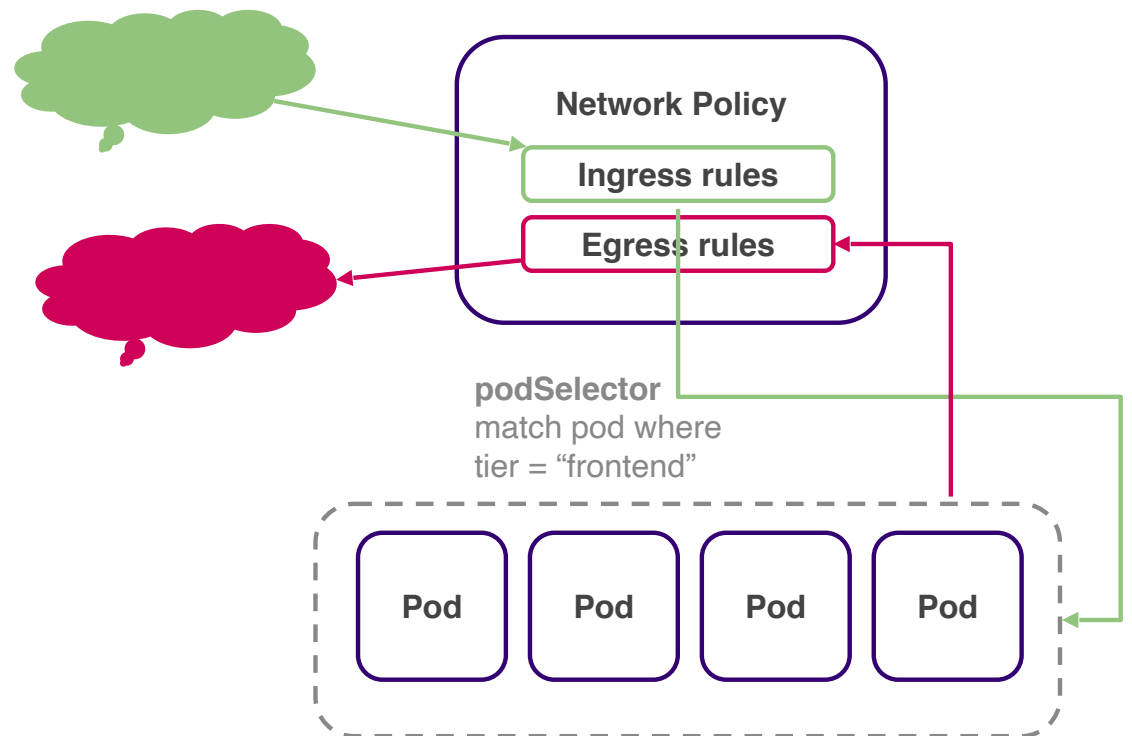
k expose pod nginx --port 8080 --type="NodePort"

# Network Policies

Kubernetes has a flat network where by default any pod can communicate with every other pod through an IP address.

Network policies are implemented using labels much like Services.

Ingress and egress rules are generally defined by the Cluster administrator. Developers can consume the policies defined by the administrator and solely focus on the application.





## Network Policies (netpol)

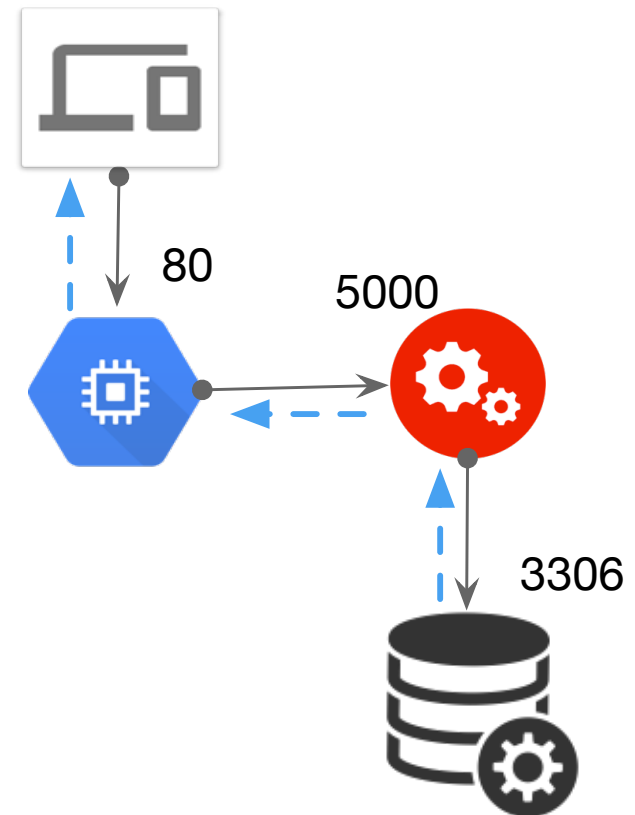
- Ingress / Egress is from components perspective

- Ingress

- Web: 80
- Mid: 5000

- Egress

- Web: 5000
- Mid: 3306



## Network Policies: Rules

- Web

Ingress: 80

Egress: 5000

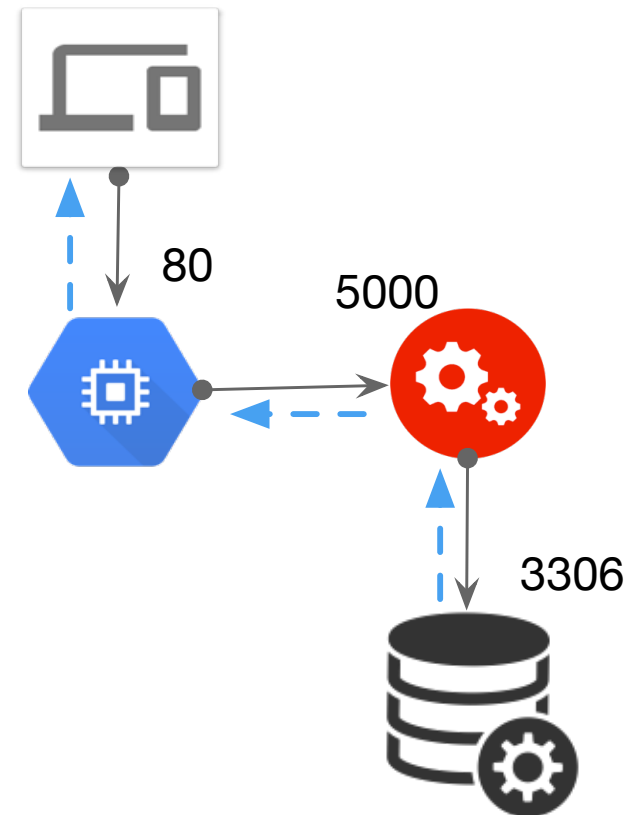
- Mid

Ingress: 5000

Egress: 3306

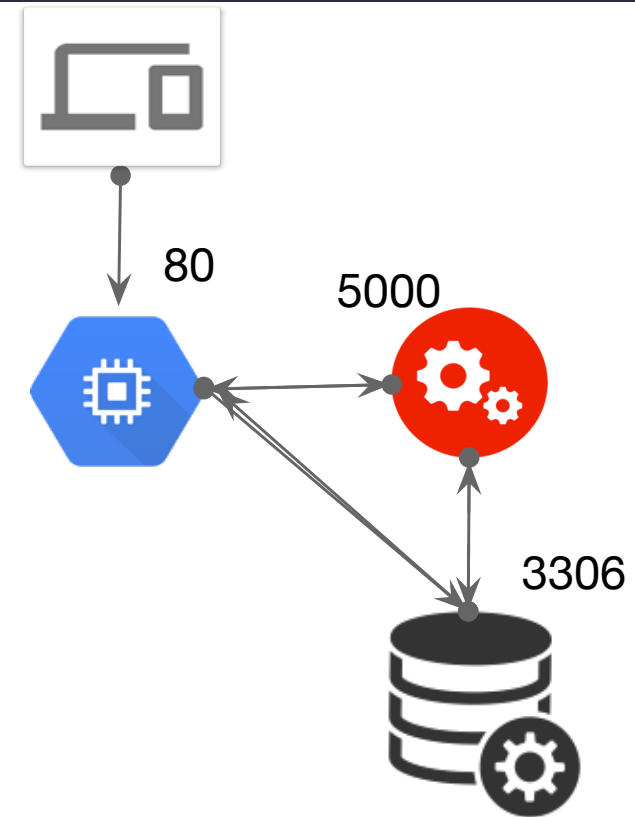
- DB

Ingress: 3306

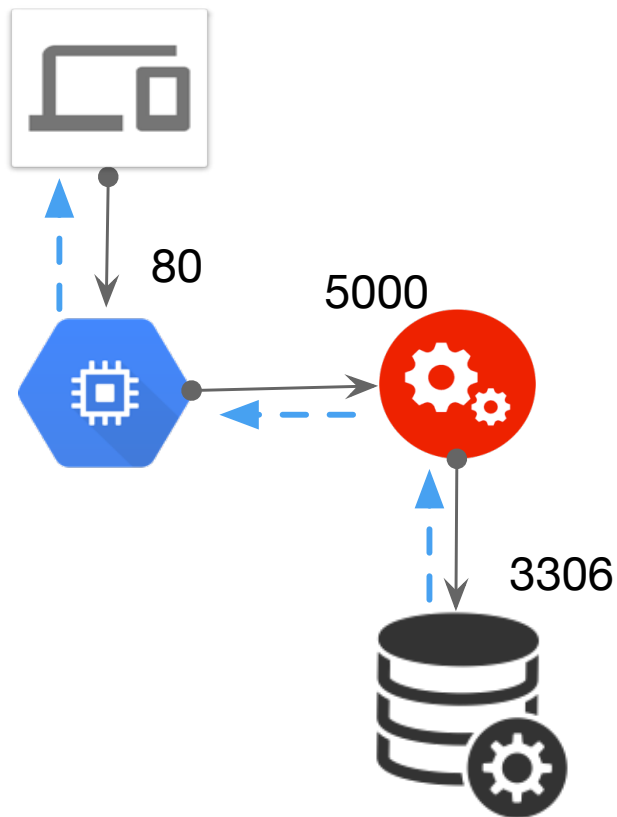


## K8S Routing Rules

- All Pods can talk to all pods



# NetworkPolicy



```
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
  name: db-policy
spec:
  podSelector:
    matchLabels:
      role: db
  policyTypes:
  - Ingress
  ingress:
  - from:
    - podSelector:
        matchLabels:
          name: api-pod
    ports:
    - protocol: TCP
      port: 3306
```

---

## NetPol Note!

Solutions support netpol

- Kube-router
- Calico
- weave-net

Do NOT support netpol

- Flannel



---

State  
Persistence  
8%

---

## State Persistence

- PV / PVC
- File Based Secrets

\* Misleading 8% - PV needed for File Based Secrets

---

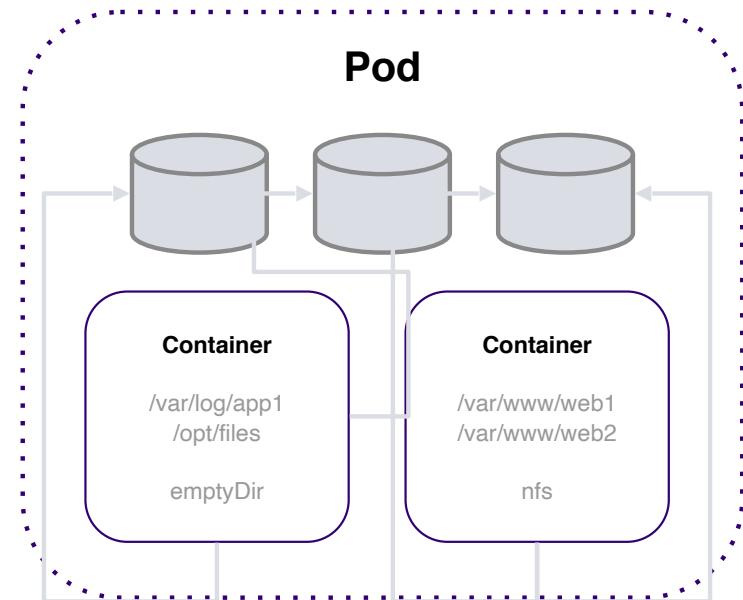
# Kubernetes Pods

- Designed to be transient
- Need a way to manage data
  - Volume mount
  - Needs various options



# Volumes

- Storage by default is ephemeral within containers
- Pods shared resources such as storage and networking
- Storage is defined as volumes and those volumes can be mounted within containers
- The underlying storage can be abstracted away through Volume types
- There are several Volume types:
  - emptyDir
  - nfs
  - hostPath
  - configMap
  - secret



Volumes definition  
`pod.spec.volumes`

Volume mounts definition  
`pod.spec.containers.volumeMounts`

## Local Volume

- Mounts local directory on the node to a path in the pod
- Needs to be mapped to a directory in a container...

```
volumes:  
  - name: data-volume  
    hostPath:  
      path: /data  
      type: Directory
```

## Local Volume

- The node path “/data” will be mapped into the container “webapp” to a mount point “/opt”
- The mapping for the mount point is via the name “data-volume”
- Data written to the container to /opt will be on the host /data

```
apiVersion: v1
kind: Pod
metadata:
  name: webapp
  labels:
    name: webapp
spec:
  containers:
  - name: webapp
    image: nginx
    ports:
      - containerPort: 80
    volumeMounts:
      - mountPath: /opt
        name: data-volume
  volumes:
  - name: data-volume
    hostPath:
      path: /data
      type: Directory
```

---

## Volume Storage Options

- Kubernetes supports lots of underlying solutions
  - NFS
  - GlusterFS
  - Ceph
  - Flocker
  - Various Cloud solutions (AWS, GCP, Azure)

## Volume Storage Options

- Kubernetes supports lots of underlying solutions
  - NFS
  - GlusterFS
  - Ceph
  - Flocker
  - Various Cloud solutions (AWS, GCP, Azure)

### AWS EBS

```
volumes:  
  - name: data-volume  
    awsElasticBlockStore:  
      volumeID: <vol-id>  
      fsType: ext4
```

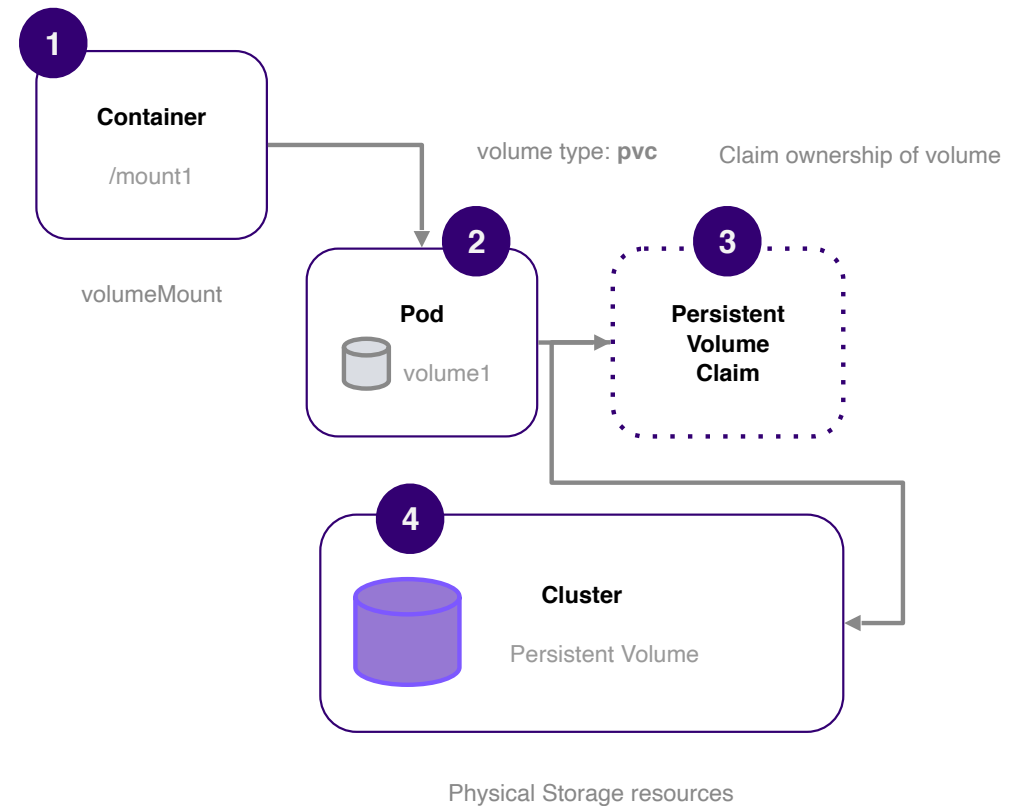
---

## Persistent Volumes

- Provides ability to centrally manage a “pool” of data storage volumes
- Pods can place a claim (called Persistent Volume Claims (PVC)) on a volume

# Persistent Volumes

- Provides ability to centrally manage a “pool” of data storage volumes
- Pods can place a claim
  - Persistent Volume Claims (PVC)



# Persistent Volume

- Access Modes:
  - ReadOnlyMany
  - ReadWriteOnce
  - ReadWriteMany

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv-vol1
spec:
  accessModes:
    - ReadWriteOnce
```



# Persistent Volume

```
kensipe@kens-mbp-2: ~/presentations/2020/ckad/lab4
k create -f lab14-2.yaml
persistentvolume/pv-vol1 created
k get pv
NAME          CAPACITY  ACCESS MODES  RECLAIM POLICY  STATUS    CLAIM  STORAGECLASS  REASON  AGE
pv-vol1       1Gi       RWO           Retain          Available  <none>  <default>     <none>  10s
```

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv-vol1
spec:
  accessModes:
    - ReadWriteOnce
  capacity:
    storage: 1Gi
  hostPath:
    path: /tmp/data
```

---

## PV vs PVC

- Admin creates PV and manages storage
- User creates a claim on that PV or something that closes matches the requirements for the pods storage needs.
- PVC are Bound to a PV based on request properties
- It is possible to bind to a volume by selectors
- Once a PV is claimed.. it can't be bound by another claim
- If a PVC can't match a PV... it remains in a Pending State.

PVC

```
kens@kens-mbp-2: ~/presentations/2020/ckad/labs
k create -f lab18-2.yaml
persistentvolume/pv-vol1 created
k get pv
NAME          CAPACITY  ACCESS MODES  RECLAIM POLICY  STATUS    CLAIM  STORAGECLASS  REASON  AGE
pv-vol1       3Gi       RWO           Retain          Available    
k get pvc
No resources found in default namespace.
k create -f lab18-3.yaml
persistentvolumeclaim/pv-claim created
k get pvc
NAME          STATUS    VOLUME          CAPACITY  ACCESS MODES  STORAGECLASS  AGE
pv-claim      Pending     
➤ ➤ ➤ ~/presentations/2020/ckad/labs
```

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: pv-claim
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 500Mi
```

---

## PVC Retention

- persistentVolumeReclaimPolicy:
  - Retain
  - Delete
  - Recycle

# Pod using PVC

```
kensipe@kens-mbp-2: ~/presentations/2020/ckad/labs
❯ k create -f lab10-2.yaml
persistentvolume/pv-vol1 created
❯ k get pv
NAME      CAPACITY  ACCESS MODES  RECLAIM POLICY  STATUS    CLAIM  STORAGECLASS  REASON  AGE
pv-vol1   500Mi     RWO           Retain          Available  pv-claim  standard      4s
❯ k create -f lab10-3.yaml
persistentvolumeclaim/pv-claim created
❯ k get pvc
NAME      STATUS    VOLUME      CAPACITY  ACCESS MODES  STORAGECLASS  AGE
pv-claim  Pending   pv-claim     500Mi     RWO           standard      3s
❯ k get pvc
NAME      STATUS    VOLUME      CAPACITY  ACCESS MODES  STORAGECLASS  AGE
pv-claim  Pending   pv-claim     500Mi     RWO           standard      27s
❯ k event
Error: unknown command "event" for "kubectl"
Run 'kubectl --help' for usage.
❯ k get events
LAST SEEN   TYPE      REASON              OBJECT                                          MESSAGE
119s       Normal    WaitForFirstConsumer persistentvolumeclaim/pv-claim                 waiting for first consumer to be created before binding
9s         Normal    WaitForFirstConsumer persistentvolumeclaim/pv-claim                 waiting for first consumer to be created before binding
11s        Normal    Killing            pod/webapp-deploy-7f4cd889-pr2rw             Stopping container nginx-container
11s        Normal    Killing            pod/webapp-deploy-7f4cd889-sxjgz             Stopping container nginx-container
11s        Normal    Killing            pod/webapp-deploy-7f4cd889-wtx9s             Stopping container nginx-container
11s        Normal    Scheduled          pod/webapp                                     Successfully assigned default/webapp to kind-worker
11s        Normal    Pulling            pod/webapp                                     Pulling image "nginx"
11s        Normal    Pulled             pod/webapp                                     Successfully pulled image "nginx"
11s        Normal    Created            pod/webapp                                     Created container webapp
11s        Normal    Started            pod/webapp                                     Started container webapp
❯ k create -f lab10-4.yaml
Error from server (AlreadyExists): error when creating "lab10-4.yaml": pods "webapp" already exists
❯ k delete pod webapp
pod "webapp" deleted
❯ k create -f lab10-4.yaml
pod/webapp created
❯ k get pvc
NAME      STATUS    VOLUME      CAPACITY  ACCESS MODES  STORAGECLASS  AGE
pv-claim  Bound     pvc-a20f8a5e-37da-432d-8d16-d0a884bf1389    500Mi     RWO           standard      5s19s
```

apiVersion: v1

kind: Pod

metadata:

name: webapp

labels:

name: webapp

spec:

containers:

- name: webapp

image: nginx

ports:

- containerPort: 80

volumeMounts:

- mountPath: /var/www/html

name: mystate

volumes:

- name: mystate

persistentVolumeClaim:

claimName: pv-claim

## Pod using PVC

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: pv-vol1
spec:
  accessModes:
    - ReadWriteOnce
  capacity:
    storage: 1Gi
  hostPath:
    path: /tmp/data
```

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: pv-claim
spec:
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 500Mi
```

```
apiVersion: v1
kind: Pod
metadata:
  name: webapp
  labels:
    name: webapp
spec:
  containers:
    - name: webapp
      image: nginx
      ports:
        - containerPort: 80
      volumeMounts:
        - mountPath: /var/www/html
          name: mystate
  volumes:
    - name: mystate
      persistentVolumeClaim:
        claimName: pv-claim
```

# Storage Classes


- Mechanism for auto-creating storage (PV) on demand

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: google-storage
provisioner: kubernetes.io/gce-pd
```

## PVC with Storage Class

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: google-storage
provisioner: kubernetes.io/gce-pd
```

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: pv-claim
spec:
  accessModes:
    - ReadWriteOnce
  storageClassName: google-storage
  resources:
    requests:
      storage: 500Mi
```





## StatefulSet

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: webapp-rs
  labels:
    name: webapp
spec:
  template:
    metadata:
      name: webapp
      labels:
        app: front-end
    spec:
      containers:
        - name: nginx-container
          image: nginx
  replicas: 3
  selector:
    matchLabels:
      app: front-end
```



```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: webapp-deploy
  labels:
    name: webapp
spec:
  template:
    metadata:
      name: webapp
      labels:
        app: front-end
    spec:
      containers:
        - name: nginx-container
          image: nginx
  replicas: 3
  selector:
    matchLabels:
      app: front-end
  serviceName: webapp-h
```

# StatefulSets

- When scaling up (or down)
  - Ordered
  - Graceful Deployment
  - Stable, unique DNS record
- Pod will only start if the previous pod is “ready”

mysql-0   mysql-1   mysql-2

---

## Headless Services

- Standard “service” load-balances across pods
- A common data source trait, is that reads can originate from all nodes, but writes need to be coordinated or be written by a “master”
- When need a DNS entry for access to an instance of a service without load balancing across the service nodes.

## Service Part of Headless Service

```
apiVersion: v1
kind: Service
metadata:
  name: webapp-h
spec:
  ports:
    - port: 80
  selector:
    app: front-end
  clusterIP: None
```

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: webapp-deploy
  labels:
    name: webapp
spec:
  template:
    metadata:
      name: webapp
      labels:
        app: front-end
    spec:
      containers:
        - name: nginx-container
          image: nginx
      replicas: 3
      selector:
        matchLabels:
          app: front-end
      serviceName: webapp-h
```

All Rights Reserved.

## Headless DNS

webapp-0.webapp-h.dev.svc.cluster.local

stateful pod name

service name

namespace service

domain

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: webapp-deploy
  labels:
    name: webapp
spec:
  template:
    metadata:
      name: webapp
      labels:
        app: front-end
    spec:
      containers:
        - name: nginx-cont
          image: nginx
  replicas: 3
  selector:
    matchLabels:
      app: front-end
  serviceName: webapp-h
```

**Thank You!**

**@kensipe**  
**kensipe@gmail.com**